


SVM, Bayesian Networks, and Ensemble Methods

SVM: Support Vector Machines

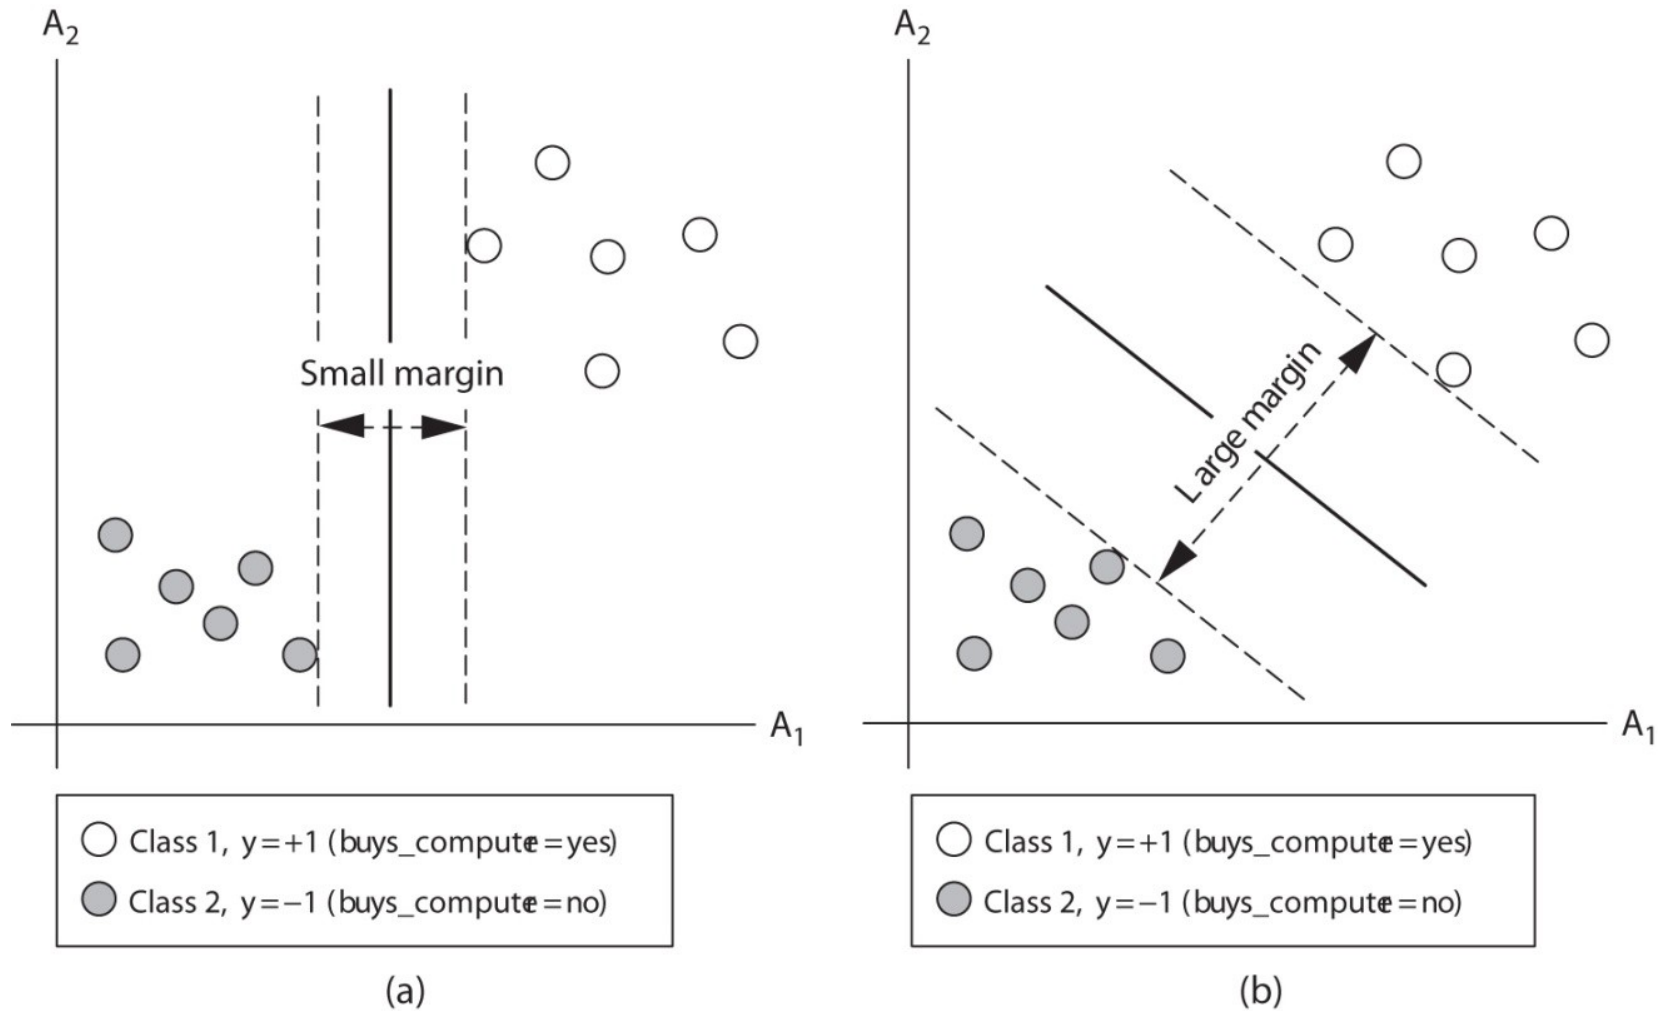
- Uses a **nonlinear** mapping to transform the original training data into a **higher** dimension
- With the new dimension, it searches for the **linear** optimal separating hyperplane (i.e., “decision boundary”)
- With an appropriate nonlinear mapping to a sufficiently high dimension, data from two classes can always be separated by a hyperplane
- SVM finds this hyperplane using support vectors (“essential” training tuples) and margins (defined by the support vectors)

Which Hyperplane?

- SVMs *maximize the margin* around the separating hyperplane
- The decision function is fully specified by a subset of training samples (support vectors)
- Quadratic programming problem



SVM: General Philosophy



Maximum Marginal Hyperplane

Why is SVM Effective on *High Dimensional* Data?

- The complexity of trained classifier is characterized by the **# of support vectors** rather than the dimensionality of the data
- The support vectors are the essential on critical training examples —they lie closest to the decision boundary (MMH)
- ❖ If all other training examples are removed and the training is repeated, the **same** separating hyperplane would be found
- The number of support vectors found can be used to compute an (upper) bound on the expected error rate of the SVM classifier, which is independent of the data dimensionality
- Thus, an SVM with a small number of support vectors can have good generalization, even when the dimensionality of the data is high

SVM: When Data is Linearly Separable

- Let data D be $(X_1, y_1), \dots, (X_{|D|}, y_{|D|})$, where X_i is the set of training tuples associated with the class labels y_i
- There are *infinite* lines (hyperplanes) separating the two classes but we want to find the **best** one (the one that *minimizes classification error on unseen data*)
- SVM searches for the hyperplane with the **largest** margin, i.e., **maximum marginal hyperplane** (MMH)

SVM: Linearly Separable

- A separating hyperplane can be written as

- $\mathbf{W} \bullet \mathbf{X} + b = 0$

where $\mathbf{W} = \{w_1, w_2, \dots, w_n\}$ is a weight vector and b a scalar (bias)

- For 2-D it can be written as

- $w_0 + w_1 x_1 + w_2 x_2 = 0$

where w_0 is the bias

- Thus any point lying above the separating hyperplane satisfies:

- $w_0 + w_1 x_1 + w_2 x_2 > 0$

- Any point lying below the separating hyperplane satisfies:

- $w_0 + w_1 x_1 + w_2 x_2 < 0$


- Thus the hyperplane defining the sides of the margin can be written as:

- $H1: w_0 + w_1 x_1 + w_2 x_2 \geq 1$ for $y_i = +1$, and

- $H2: w_0 + w_1 x_1 + w_2 x_2 \leq -1$ for $y_i = -1$

- Any training tuples falling on hyperplanes $H1$ or $H2$ (i.e., the sides defining the margin) are support vectors
- This becomes a **constrained (convex) quadratic optimization problem**: Quadratic objective function and linear constraints; Quadratic Programming (QP); Lagrangian multipliers

Linear SVM Mathematically



What we know:

- $w \cdot x^+ + b = +1$
- $w \cdot x^- + b = -1$
- $w \cdot (x^+ - x^-) = 2$

$$M = \frac{(x^+ - x^-) \cdot w}{|w|} = \frac{2}{|w|}$$

Linear SVM Mathematically

- Goal: 1) Correctly classify all training data

$$wx_i + b \geq 1 \quad \text{if } y_i = +1$$


$$wx_i + b \leq -1 \quad \text{if } y_i = -1$$

$$y_i(wx_i + b) \geq 1 \quad \text{for all } i$$

- 2) Maximize the Margin

$$M = \frac{2}{|w|}$$

$$\text{same as minimize } \frac{1}{2} w^t w$$



- We can formulate a Quadratic Optimization Problem and solve for w and b

Minimize	$\Phi(w) = \frac{1}{2} w^t w$
subject to	$y_i(wx_i + b) \geq 1 \quad \forall i$

Solving the Optimization Problem

Find \mathbf{w} and b such that
 $\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$ is minimized;
and for all $\{(\mathbf{x}_i, y_i)\}$: $y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$

- Need to optimize a quadratic function subject to linear constraints.
- Quadratic optimization problems are a well-known class of mathematical programming problems, and many (rather intricate) algorithms exist for solving them.
- The solution involves constructing a *dual problem* where a *Lagrange multiplier* α_i is associated with every constraint in the primary problem:

Find $\alpha_1 \dots \alpha_N$ such that
 $Q(\boldsymbol{\alpha}) = \sum \alpha_i - \frac{1}{2} \sum \sum \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$ is maximized and
(1) $\sum \alpha_i y_i = 0$
(2) $\alpha_i \geq 0$ for all α_i

The Optimization Problem Solution

- The solution has the form:


$$\mathbf{w} = \sum \alpha_i y_i \mathbf{x}_i \quad b = y_k - \mathbf{w}^T \mathbf{x}_k \text{ for any } \mathbf{x}_k \text{ such that } \alpha_k \neq 0$$

- Each non-zero α_i indicates that corresponding \mathbf{x}_i is a support vector.
- Then the classifying function will have the form:

$$f(\mathbf{x}) = \sum \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b$$

- Notice that it relies on an *inner product* between the test point \mathbf{x} and the support vectors \mathbf{x}_i .
- Also keep in mind that solving the optimization problem involves computing the inner products $\mathbf{x}_i^T \mathbf{x}_j$ between **all** pairs of training points.

Dataset with noise




- Hard Margin: So far we require **all** data points be classified correctly
 - - No training error
- What if the training set is noisy?
- - Solution 1: use very powerful kernels

OVERFITTING!

Soft Margin Classification

Slack variables ξ_i can be added to allow misclassification of difficult or noisy examples.



What should our quadratic optimization criterion be?

Minimize $\frac{1}{2} \mathbf{w} \cdot \mathbf{w} + C \sum_{k=1}^R \xi_k$

Hard Margin v.s. Soft Margin

- The old formulation:

Find \mathbf{w} and b such that

$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$ is minimized and for all $\{(\mathbf{x}_i, y_i)\}$

$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$$

- The new formulation incorporating slack variables:

Find \mathbf{w} and b such that

$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum \xi_i$ is minimized and for all $\{(\mathbf{x}_i, y_i)\}$

$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i \quad \text{and} \quad \xi_i \geq 0 \text{ for all } i$$

- Parameter C can be viewed as a way to control overfitting.

Linear SVMs: Summary

- The classifier is a separating hyperplane.
- Most “important” training points are support vectors; they define the hyperplane.
- Quadratic optimization algorithms can identify which training points x_i are support vectors with non-zero Lagrangian multipliers α_i .
- Both in the dual formulation of the problem and in the solution training points appear only inside dot products:

Find $\alpha_1 \dots \alpha_N$ such that

$Q(\alpha) = \sum \alpha_i - \frac{1}{2} \sum \sum \alpha_i \alpha_j y_i y_j x_i^T x_j$ is maximized and


(1) $\sum \alpha_i y_i = 0$

(2) $0 \leq \alpha_i \leq C$ for all α_i


$$f(x) = \sum \alpha_i y_i x_i^T x + b$$

SVM: Linearly **In**separable

- Extend linear SVMs
- Find non-linear decision boundaries (**non-linear** hypersurfaces).
- **Step 1**: transform the original input data into high dimensional space
- **Step 2**: Search for linear separating hyperplane in the new space
- The linear hyperplane in high dimensional space is non-linear in the original space



SVM Linearly Inseparable



SVM: Linearly Inseparable

- Requires a **mapping** to higher dimensional space
 - e.g. $\varphi([x_i, x_j]) = [x_i, x_j, x_i^2 + x_j^2]$
 - or $\varphi(X_i, X_j) = \varphi(X_i) \bullet \varphi(X_j)$
- Two problems:
 - How to choose a mapping?
 - Higher number of dimensions increases computational cost
 - Computing the dot product for all of the support vectors
 - One multiplication and one addition for each dimension
- Math trick:
 - Apply a kernel function to the original input data
 - Wherever a dot product is used, replace it with a kernel function

SVM: Kernel Function

- Equivalent to the dot product
 - $K(X_i, X_j) = \phi(X_i) \cdot \phi(X_j)$
- We can solve $K(X_i, X_j)$ without having to perform any of the mappings $\phi(X_i)$
- Provide a way to manipulate data as though **it were projected into a higher dimensional space** by operating on it in its original space

SVM: Kernel Function

- Instead of computing the dot product on the transformed data, it is mathematically equivalent to apply a kernel function $K(X_i, X_j)$ to the original data, i.e., $K(X_i, X_j) = \Phi(X_i) \cdot \Phi(X_j)$
- Typical Kernel Functions

Polynomial kernel of degree h : $K(X_i, X_j) = (X_i \cdot X_j + 1)^h$

Gaussian radial basis function kernel : $K(X_i, X_j) = e^{-\|X_i - X_j\|^2 / 2\sigma^2}$

Sigmoid kernel : $K(X_i, X_j) = \tanh(\kappa X_i \cdot X_j - \delta)$

- After applying the kernel function, finding **maximal margin hyperplane** is similar to linear SVM with an upper bound on the Lagrangian multipliers
- SVM can also be used for classifying multiple (> 2) classes and for regression analysis (with additional parameters)

Weakness of SVM

- It is **sensitive** to noise
 - A relatively small number of mislabeled examples can dramatically decrease the performance
- It only considers two classes
 - how to do multi-class classification with SVM?
 - Answer:
 - 1) with output arity **m**, learn **m** SVM's
 - SVM 1 learns "Output == 1" vs "Output != 1"
 - SVM 2 learns "Output == 2" vs "Output != 2"
 - :
 - SVM m learns "Output == m" vs "Output != m"
 - 2) To predict the output for a new input, just predict with each SVM and find out which one puts the prediction the ***furthest*** into the positive region.

SVM **vs.** Logistic Regression

- Essentially they have the *same objective*
 - Binary classification with a **linear** decision boundary
 - Both are complex optimization problems solved using gradient descent and Lagrangian methods
- Key differences
 - LR uses a logistic (sigmoid) function to classify the data
 - SVM uses the “kernel trick” to find the maximum marginal boundary (defined by support vectors)
 - “Kernel trick” is difficult to implement using LR
 - SVM has been shown to **outperform** LR

SVM Applications

- SVM has been used successfully in many real-world problems
 - text (and hypertext) categorization
 - image classification
 - bioinformatics (protein classification, cancer classification)
 - hand-written character recognition

Application 1: Cancer Classification

- High Dimensional

- $p > 1000$; $n < 100$

- Imbalanced

- less positive samples

$$K[x, x] = k(x, x) + \lambda \frac{n^+}{N}$$

- Many irrelevant features

- Noisy

SVM is sensitive to noisy (mis-labeled) data ☹

Genes				
Patients	g-1	g-2	g-p
P-1				
p-2				
.....				
p-n				

FEATURE SELECTION

In the linear case,
 w_i^2 gives the ranking of dim i

Application 2: Text Categorization

- Task: The classification of natural text (or hypertext) documents into a fixed number of predefined categories based on their content.
 - email filtering, web searching, sorting documents by topic, etc..
- A document can be assigned to more than one category, so this can be viewed as a series of binary classification problems, one for each category

Representation of Text

IR's vector space model (aka bag-of-words representation)

- A doc is represented by a vector indexed by a pre-fixed set or dictionary of terms
- Values of an entry can be binary or weights

$$\phi_i(x) = \frac{\text{tf}_i \log(\text{idf}_i)}{\kappa},$$

- Normalization, stop words, word stems
- Doc $x \Rightarrow \boldsymbol{\varphi}(x)$
- <https://www.csee.umbc.edu/~ian/irF02/lectures/07Models-VSM.pdf>

Text Categorization using SVM


- The distance between two documents is $\varphi(x) \cdot \varphi(z)$
- $K(x,z) = \varphi(x) \cdot \varphi(z)$ is a valid kernel, SVM can be used with $K(x,z)$ for discrimination.
- Why SVM?
 - High dimensional input space
 - Few irrelevant features (dense concept)
 - Sparse document vectors (sparse instances)
 - Text categorization problems are linearly separable

Some Issues in SVM

- Choice of kernel
 - Gaussian or polynomial kernel is default
 - If ineffective, more elaborate kernels are needed
 - Domain experts can give assistance in formulating appropriate similarity measures
- Choice of kernel parameters
 - e.g. σ in Gaussian kernel
 - σ is the distance between closest points with different classifications
 - In the absence of reliable criteria, applications rely on the use of a validation set or cross-validation to set such parameters.
- Optimization criterion – Hard margin v.s. Soft margin
 - a lengthy series of experiments in which various parameters are tested

Bayesian Belief Networks

- **Bayesian belief networks** (also known as **Bayesian networks**, **probabilistic networks**)
 - A simple, graphical notation for *conditional independence assertions* and hence for compact specification of full joint distributions
 - Allow *class conditional independencies* between *subsets* of variables
- A (*directed acyclic*) graphical model of causal relationships
 - Represents dependency among the variables
 - Gives a specification of **joint** probability distribution



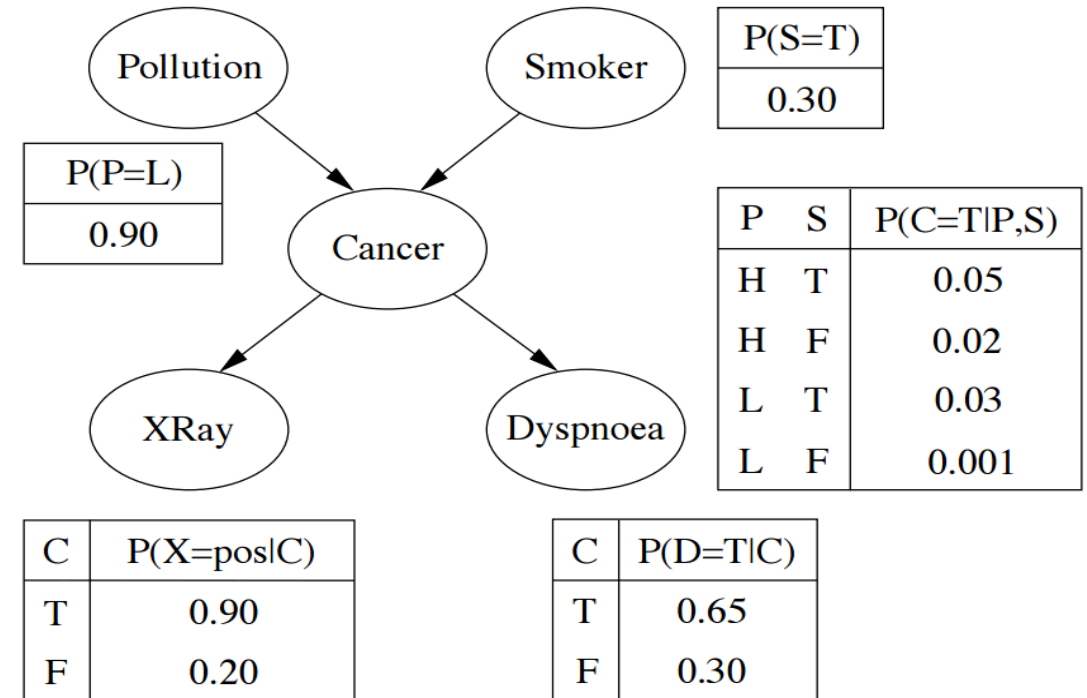
- Nodes: random variables
- Links: dependency
- X and Y are the parents of Z, and Y is the parent of P
- No dependency between Z and P
- Has no loops/cycles

Bayesian Belief Network: An Example

- Derivation of the conditional probability of a particular combination of values of \mathbf{X} , from conditional probability table (CPT):

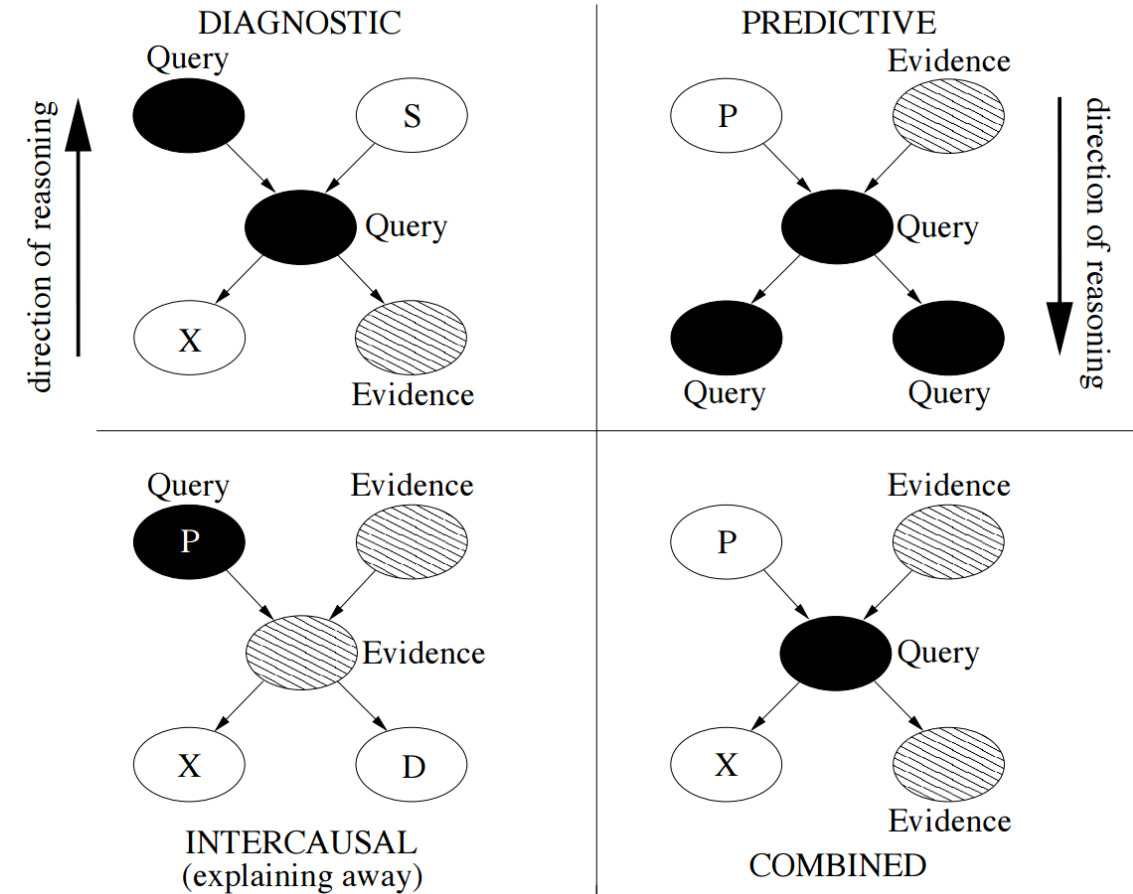
$$P(x_1, x_2, \dots, x_n) = \prod_i P(x_i | \text{Parents}(X_i))$$

$$\begin{aligned}
 &P(X = \text{pos} \wedge D = T \wedge C = T \wedge P = \text{low} \wedge S = F) \\
 &= P(X = \text{pos} | D = T, C = T, P = \text{low}, S = F) \\
 &\quad \times P(D = T | C = T, P = \text{low}, S = F) \\
 &\quad \times P(C = T | P = \text{low}, S = F) P(P = \text{low} | S = F) P(S = F) \\
 &= P(X = \text{pos} | C = T) P(D = T | C = T) P(C = T | P = \text{low}, S = F) \\
 &\quad \times P(P = \text{low}) P(S = F)
 \end{aligned}$$



Bayesian Belief Network: Types of Reasoning

- Bayesian Networks can support any direction of reasoning
- Diagnostic reasoning
 - Reasoning **from symptom to cause**
- Predictive reasoning
 - Reasoning **from new information about causes to new beliefs about effects**
- Inter-causal reasoning
 - Reasoning about the **mutual causes of a common effect**
- Combined reasoning
 - **Combining above types of reasoning** in any way



Updating Beliefs Given New Information

Node P(S)=0.3	No Evidence	Reasoning Case				
		Diagnostic D=T	Predictive S=T	Intercausal C=T C=T S=T		Combined D=T S=T
Bel(P=high)	0.100	0.102	0.100	0.249	0.156	0.102
Bel(S=T)	0.300	0.307	1	0.825	1	1
Bel(C=T)	0.011	0.025	0.032	1	1	0.067
Bel(X=pos)	0.208	0.217	0.222	0.900	0.900	0.247
Bel(D=T)	0.304	1	0.311	0.650	0.650	1
P(S)=0.5						
Bel(P=high)	0.100	0.102	0.100	0.201	0.156	0.102
Bel(S=T)	0.500	0.508	1	0.917	1	1
Bel(C=T)	0.174	0.037	0.032	1	1	0.067
Bel(X=pos)	0.212	0.226	0.311	0.900	0.900	0.247
Bel(D=T)	0.306	1	0.222	0.650	0.650	1

How are Bayesian Networks Constructed?

- **Subjective construction:** Identification of (direct) causal structure
 - People are quite good at identifying direct causes from a given set of variables & whether the set contains all relevant direct causes
 - **Markovian** assumption: Each variable becomes independent of its non-effects once its direct causes are known
 - E.g., $S \leftarrow F \rightarrow A \leftarrow T$, path $S \rightarrow A$ is blocked once we know $F \rightarrow A$
 - HMM (Hidden Markov Model): often used to model dynamic systems whose states are not observable, yet their outputs are
- **Synthesis from other specifications**
 - E.g., from a formal system design: block diagrams & info flow
- **Learning from data**
 - Learning parameters gives its structure or learning both structure and parameters
 - Maximum likelihood principle: favors Bayesian networks that maximize the probability of observing the given data set

Training Bayesian Networks

- **Scenario 1:** Given both the network structure and all variables observable: *compute only the CPT entries*
- **Scenario 2:** Network structure known, some variables hidden (missing values): ***gradient descent*** method, i.e., search for a solution along the steepest descent of a criterion function
 - Weights are initialized to random probability values
 - At each iteration, it moves towards what appears to be the best solution at the moment, without backtracking
 - Weights are updated at each iteration & converge to local optimum
- **Scenario 3:** Network structure unknown, all variables observable: search through the model space to *reconstruct network topology*
- **Scenario 4:** Unknown structure, all hidden variables: No good algorithms known for this purpose
- More:
 - D. Heckerman. A Tutorial on Learning with Bayesian Networks. In *Learning in Graphical Models*, M. Jordan, ed.. MIT Press, 1999.

Ensemble Methods: Increasing Accuracy

- Ensemble methods
 - Use a **combination** of models to **increase** accuracy
 - Combine a series of k learned models, M_1, M_2, \dots, M_k , with the aim of creating an **improved** model M^*
- Popular ensemble methods
 - **Bagging**: averaging the prediction over a collection of classifiers
 - **Boosting**: weighted vote with a collection of classifiers
 - **Random Forests**
 - **Ensemble**: combining a set of heterogeneous classifiers

Bagging: Bootstrap Aggregation

- **Analogy:** Diagnosis based on multiple doctors' majority vote
- Training
 - Given a set D of d tuples, at each iteration i , a training set D_i of d tuples is sampled with replacement from D (i.e., bootstrap)
 - A classifier model M_i is learned for each training set D_i
- Classification: classify an unknown sample X
 - Each classifier M_i returns its class prediction
 - The bagged classifier M^* counts the votes and assigns the class with the **most votes** to X
- Prediction: can be applied to the prediction of continuous values by taking the average value of each prediction for a given test tuple
- Accuracy
 - Often significantly better than a single classifier derived from D
 - For noise data: not considerably worse, more **robust**
 - Proved improved accuracy in prediction

Boosting

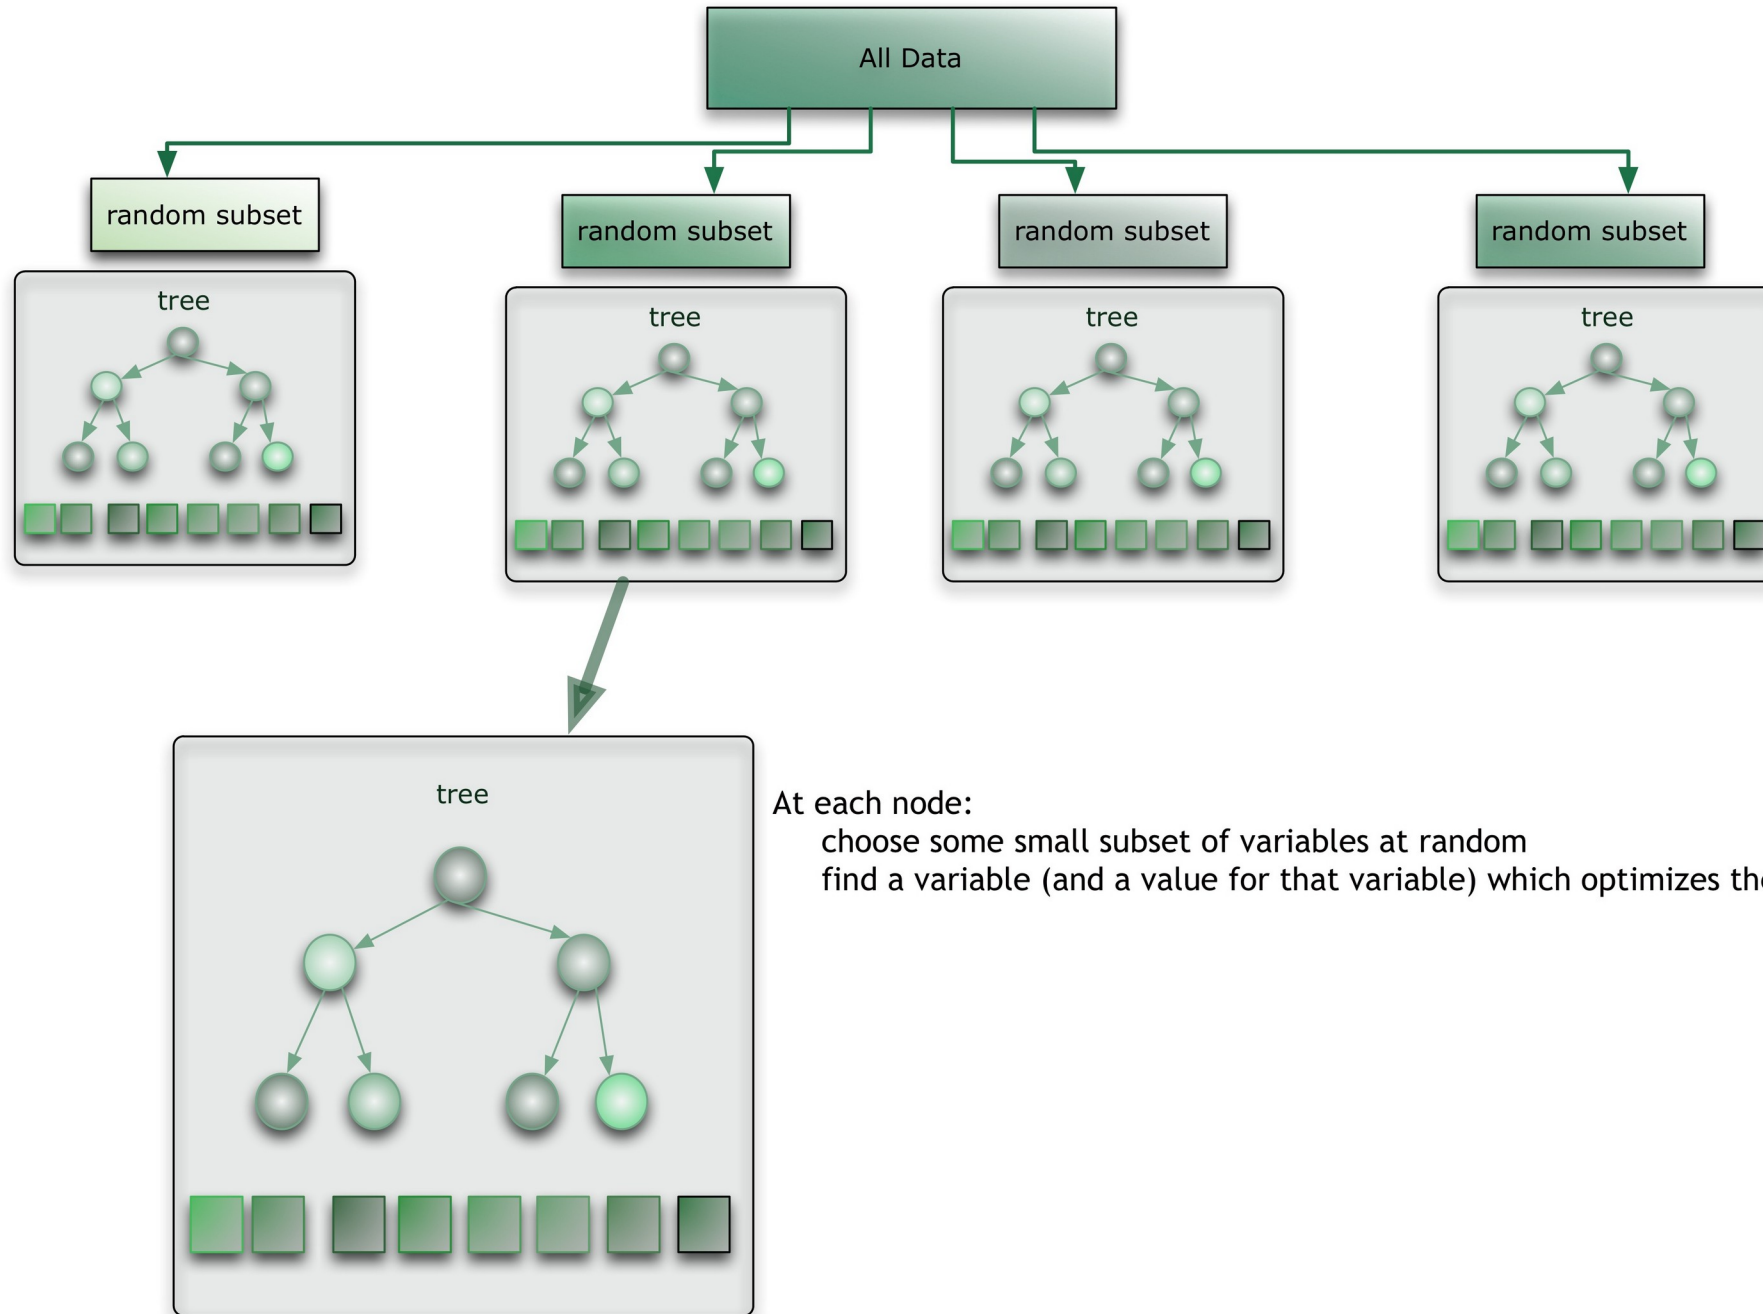
- **Analogy:** Consult several doctors, based on a combination of weighted diagnoses - weight assigned based on the previous diagnosis accuracy
- How boosting works?
 - **Weights** are assigned to each training tuple
 - A series of k classifiers is iteratively learned
 - After a classifier M_i is learned, the weights are updated to allow the subsequent classifier, M_{i+1} , to **pay more attention to the training tuples that were misclassified** by M_i
 - The final classifier M^* **combines the votes** of each individual classifier, where the weight of each classifier's vote is a function of its accuracy
- Boosting algorithm can be extended for numeric prediction
- Comparing with bagging: Boosting tends to have greater accuracy, but it also risks overfitting the model to misclassified data

Random Forest (Breiman 2001)

- Random Forest:
 - Each classifier in the ensemble is a *decision tree* classifier and is generated using a **random selection of attributes at each node** to determine the split
 - During classification, each tree votes and the most popular class is returned
- Two Methods to construct Random Forest:
 - **Forest-RI** (*random input selection*): Randomly select, at each node, F attributes as candidates for the split at the node. The CART methodology is used to grow the trees to maximum size
 - **Forest-RC** (*random linear combinations*): Creates new attributes (or features) that are a **linear** combination of the existing attributes (reduces the correlation between individual classifiers)
- Comparable in accuracy to boosting, but **more robust to errors and outliers**
- Insensitive to the number of attributes selected for consideration at each split, and faster than bagging or boosting

General Random Forest Algorithm

- Each tree is constructed using the following algorithm:
 1. N : the number of training samples, M : the number of variables in the classifier.
 2. m : the number of input variables to be used to determine the decision at a node of the tree; m should be much less than M .
 3. Choose a training set for this tree by choosing n times with replacement from all N available training cases (i.e.: bootstrap). Use the rest of the cases to estimate the error of the tree, by predicting their classes.
 4. For each node of the tree, randomly choose m variables on which to base the decision at that node. Calculate the best split based on these m variables in the training set.
 5. Each tree is fully grown and not pruned.
- Prediction: a new sample is assigned the label of the training sample in the terminal node it ends up in. This procedure is iterated over all trees in the ensemble, and the average vote of all trees is reported as random forest prediction.



References

- <https://scikit-learn.org/stable/modules/svm.html>
- <https://towardsdatascience.com/bbn-bayesian-belief-networks-how-to-build-them-effectively-in-python-6b7f93435bba>
- <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>