

Sequential Pattern Mining

What is Sequential Pattern Mining

- Find patterns in data where items are delivered in a **sequence**.
- A sequence is an ordered list of events
- Examples include:
 - Time series data
 - Symbolic sequences
 - Biological sequences

Applications

- Applications of sequential pattern mining
 - Customer shopping sequences
 - Medical treatments
 - Natural disasters (e.g., earthquakes),
 - Science & engineering processes
 - Stocks and markets, etc.
 - Telephone calling patterns, Weblog click streams
 - DNA sequences and gene structures
 - Sports data mining

The Traditional CS Example

- A sequence database consists of ordered elements or events
- Transactions are orderless (e.g. diapers and beer)
- Sequences include some **order** (e.g. diapers then beer)
- Transaction databases vs. sequence databases

A transaction database

TID	itemsets
10	a, b, d
20	a, c, d
30	a, d, e
40	b, e, f

A sequence database

SID	sequences
10	<a(abc)(ac)d(cf)>
20	<(ad)c(bc)(ae)>
30	<(ef)(ab)(df)cb>
40	<eg(af)cbc>

Subsequence vs. super sequence

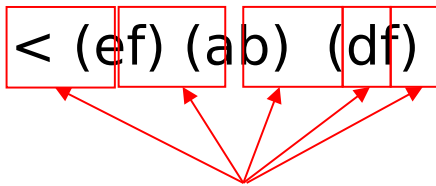
- A sequence is an ordered list of events, denoted $\langle e_1 e_2 \dots e_l \rangle$
- Given two sequences $\alpha = \langle a_1 a_2 \dots a_n \rangle$ and $\beta = \langle b_1 b_2 \dots b_m \rangle$
- α is called a **subsequence** of β , denoted as $\alpha \subseteq \beta$, if there exist integers $1 \leq j_1 < j_2 < \dots < j_n \leq m$ such that $a_1 \subseteq b_{j_1}, a_2 \subseteq b_{j_2}, \dots, a_n \subseteq b_{j_n}$
- β is a super sequence of α
 - E.g. $\alpha = \langle \text{ab}, \text{d} \rangle$ and $\beta = \langle \text{abc}, \text{de}, \text{abc} \rangle$

What Is Sequential Pattern Mining?

- Given a set of sequences and **support threshold**, find the complete set of *frequent* subsequences

A sequence database

SID	sequence
10	<a(<u>ab</u> c)(a <u>c</u>)d(cf)>
20	<(ad)c(bc)(ae)>
30	<(ef)(<u>ab</u>)(df) <u>c</u> b>
40	<eg(af)cbc>

A sequence : < (ef) (ab) (df) c b >
 >
 

An element may contain a set of items e.g. (ef). Items within an element are unordered and we list them alphabetically._

<a(bc)dc> is a subsequence of <a(abc)(ac)d(cf)>

Given support threshold $min_sup = 2$, <(ab)c> is a sequential pattern

Challenges in Sequential Pattern Mining

- A **huge** number of possible sequential patterns are hidden in databases
- A mining algorithm should
 - find the **complete set of patterns**, when possible, satisfying the minimum support (frequency) threshold
 - be highly **efficient and scalable** involving only a small number of database scans
 - be able to incorporate various kinds of **user-specific constraints**

Thought Exercise

- How can the **a-priori** property be extended to sequential patterns?
- What would an a-priori-like algorithm for sequential pattern mining look like?
 - Candidate generation?

Studies on Sequential Pattern Mining

- Concept introduction and an initial Apriori-like algorithm
 - Agrawal & Srikant. Mining sequential patterns, [ICDE'95]
- Apriori-based method: **GSP** (Generalized Sequential Patterns: Srikant & Agrawal [EDBT'96])
- Pattern-growth methods: FreeSpan & **PrefixSpan** (Han et al. KDD'00; Pei, et al. [ICDE'01])
- Vertical format-based mining: **SPADE** (Zaki [Machine Learning'00])
- Constraint-based sequential pattern mining (SPIRIT: Garofalakis, Rastogi, Shim [VLDB'99]; Pei, Han, Wang [CIKM'02])
- Mining closed sequential patterns: **CloSpan** (Yan, Han & Afshar [SDM'03])

Methods for sequential pattern mining

- **Apriori-based Approaches**
 - GSP
 - SPADE
- **Pattern-Growth-based Approaches**
 - FreeSpan
 - PrefixSpan

The Apriori Property of Sequential Patterns

- A basic property: Apriori (Agrawal & Srikant'94)
 - If a sequence S is not frequent, then **none of the super-sequences of S is frequent**
 - E.g, $\langle hb \rangle$ is infrequent so are $\langle hab \rangle$ and $\langle (ah)b \rangle$

Seq. ID	Sequence
10	$\langle (bd)cb(ac) \rangle$
20	$\langle (bf)(ce)b(fg) \rangle$
30	$\langle (ah)(bf)abf \rangle$
40	$\langle (be)(ce)d \rangle$
50	$\langle a(bd)bcb(ade) \rangle$

Given support threshold $min_sup = 2$

GSP-- Generalized Sequential Pattern Mining

- GSP (Generalized Sequential Pattern) mining algorithm
- Outline of the method
 - Initially, every item in DB is a candidate of **length-1**
 - for each level (i.e., sequences of length-k) do
 - scan database to **collect support count for each candidate sequence**
 - **generate candidate length-(k+1) sequences** from length-k frequent sequences using Apriori
 - repeat until no frequent sequence or no candidate can be found
- Major strength: Candidate pruning by **Apriori**

Finding Length-1 Sequential Patterns

- Initial candidates:
 - <a>, , <c>, <d>, <e>, <f>, <g>, <h>
- Scan database once, count support for candidates

min_sup

Seq. ID	Sequence
10	<(bd)cb(ac)>
20	<(bf)(ce)b(fg)>
30	<(ah)(bf)abf>
40	<(be)(ce)d>
50	<a(bd)bcb(ade)>

Cand	Sup
<a>	3
	5
<c>	4
<d>	3
<e>	3
<f>	2
<g>	1
<h>	1

Generating Length-2 Candidates

51 length-2
Candidates

	<a>		<c>	<d>	<e>	<f>
<a>	<aa>	<ab>	<ac>	<ad>	<ae>	<af>
	<ba>	<bb>	<bc>	<bd>	<be>	<bf>
<c>	<ca>	<cb>	<cc>	<cd>	<ce>	<cf>
<d>	<da>	<db>	<dc>	<dd>	<de>	<df>
<e>	<ea>	<eb>	<ec>	<ed>	<ee>	<ef>
<f>	<fa>	<fb>	<fc>	<fd>	<fe>	<ff>

	<a>		<c>	<d>	<e>	<f>
<a>		<(ab)>	<(ac)>	<(ad)>	<(ae)>	<(af)>
			<(bc)>	<(bd)>	<(be)>	<(bf)>
<c>				<(cd)>	<(ce)>	<(cf)>
<d>					<(de)>	<(df)>
<e>						<(ef)>
<f>						

Without Apriori property,
 $8*8+8*7/2=92$ candidates

Apriori prunes

44.57%
candidates

Finding Length-2 Sequential Patterns

- Scan **updated** database one more time, collect support count for each length-2 candidate
- There are 19 length-2 candidates which pass the minimum support threshold
 - They are length-2 sequential patterns

min_sup

Seq. ID	Sequence
10	<(bd)cb(ac)>
20	<(bf)(ce)b(fg)>
30	<(ah)(bf)abf>
40	<(be)(ce)d>
50	<a(bd)bcb(ade)>

The GSP Mining Process

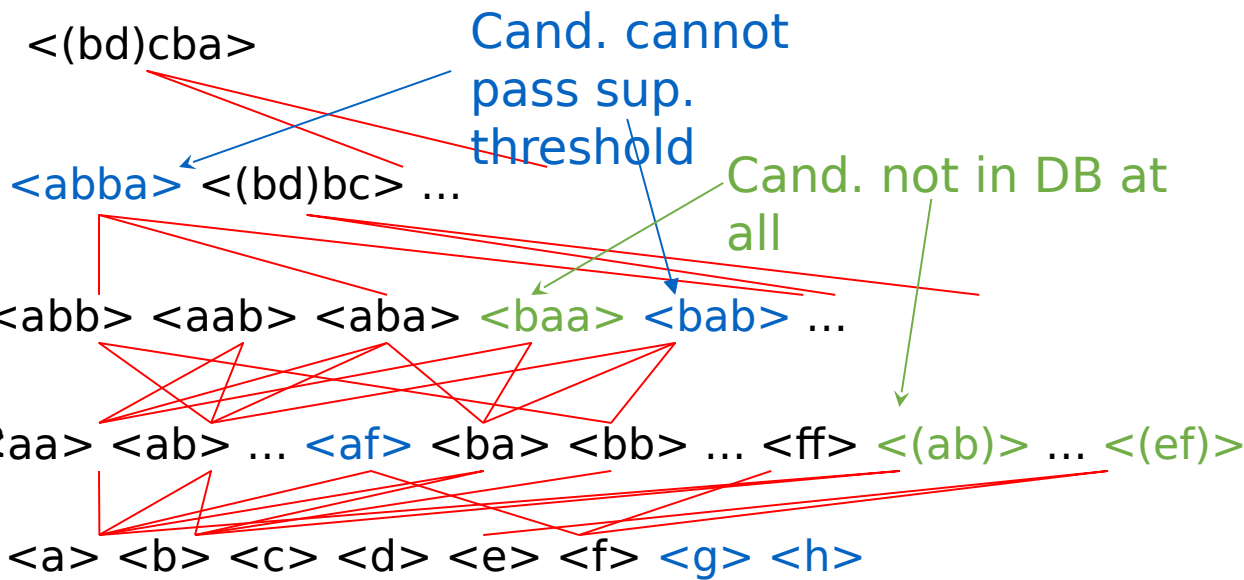
5th scan: 1 cand. 1 length-5 seq.

4th scan: 8 cand. 6 length-4 seq.

3rd scan: 46 cand. 19 length-3 seq. 20 cand. not in DB at all

2nd scan: 51 cand. 19 length-2 seq.

1st scan: 8 cand. 6 length-1 seq.



$min_sup = 2$

Seq. ID	Sequence
10	$\langle (bd)cb(ac) \rangle$
20	$\langle (bf)(ce)b(fg) \rangle$
30	$\langle (ah)(bf)abf \rangle$
40	$\langle (be)(ce)d \rangle$
50	$\langle a(bd)bcb(ade) \rangle$

The GSP Algorithm

F_1 = the set of frequent 1-sequence

$k=2$,

while $F(k-1)$ is not empty;

 Generate candidate sets C_k (set of candidate k -sequences);

 For all input sequences s in the database D

 Increment count of all a in C_k if s supports a

$F_k = \{a \in C_k \text{ such that its frequency exceeds the threshold}\}$

$k = k+1$;

 Result = Set of all frequent sequences is the union of all F_k s

End

End

The GSP Algorithm

- Benefits from Apriori pruning
 - Reduces search space
- Bottlenecks
 - Scans the database multiple times
 - Generates a huge set of candidate sequences

The SPADE Algorithm

- SPADE (Sequential Pattern Discovery using Equivalent Class) developed by Zaki 2001
- A **vertical** format sequential pattern mining method
- A sequence database is mapped to a large set of Items: <SID, EID>
- Sequential pattern mining is performed by
 - growing the subsequences (patterns) one item at a time by Apriori candidate generation

The SPADE Algorithm

SID	EID	Items
1	1	a
1	2	abc
1	3	ac
1	4	d
1	5	cf
2	1	ad
2	2	c
2	3	bc
2	4	ae
3	1	ef
3	2	ab
3	3	df
3	4	c
3	5	b
4	1	e
4	2	g
4	3	af
4	4	c
4	5	b
4	6	c

a		b		...
SID	EID	SID	EID	...
1	1	1	2	
1	2	2	3	
1	3	3	2	
2	1	3	5	
2	4	4	5	
3	2			
4	3			

ab			ba			...
SID	EID (a)	EID(b)	SID	EID (b)	EID(a)	...
1	1	2	1	2	3	
2	1	3	2	3	4	
3	2	5				
4	3	5				

aba				...
SID	EID (a)	EID(b)	EID(a)	...
1	1	2	3	
2	1	3	4	

Bottlenecks of Candidate Generate-and-test

- A huge set of candidates generated.
 - Especially 2-item candidate sequence.
- Multiple Scans of database in mining.
 - The length of each candidate grows by one at each database scan.
- Inefficient for mining long sequential patterns.
 - A long pattern grows up from short patterns
 - An exponential number of short candidates

PrefixSpan (Prefix-Projected Sequential Pattern Growth)

- PrefixSpan
 - Projection-based
 - But only prefix-based projection: less projections and quickly shrinking sequences
- J.Pei, J.Han,... PrefixSpan : Mining sequential patterns efficiently by prefix-projected pattern growth. ICDE'01.

Prefix and Suffix (Projection)

- Given sequence $\langle a(abc)(ac)d(cf) \rangle$
- $\langle a \rangle$, $\langle aa \rangle$, $\langle a(ab) \rangle$ and $\langle a(abc) \rangle$ are prefixes of sequence $\langle a(abc)(ac)d(cf) \rangle$

Prefix	<u>Suffix</u> (Prefix-Based <u>Projection</u>)
$\langle a \rangle$	$\langle (abc)(ac)d(cf) \rangle$
$\langle aa \rangle$	$\langle (_bc)(ac)d(cf) \rangle$
$\langle ab \rangle$	$\langle (_c)(ac)d(cf) \rangle$

Mining Sequential Patterns by Prefix Projections

- **Step 1:** find length-1 sequential patterns
 - $\langle a \rangle$, $\langle b \rangle$, $\langle c \rangle$, $\langle d \rangle$, $\langle e \rangle$, $\langle f \rangle$
- **Step 2:** **divide** search space. The complete set of seq. pat. can be partitioned into 6 subsets:
 - The ones having prefix $\langle a \rangle$;
 - The ones having prefix $\langle b \rangle$;
 - ...
 - The ones having prefix $\langle f \rangle$

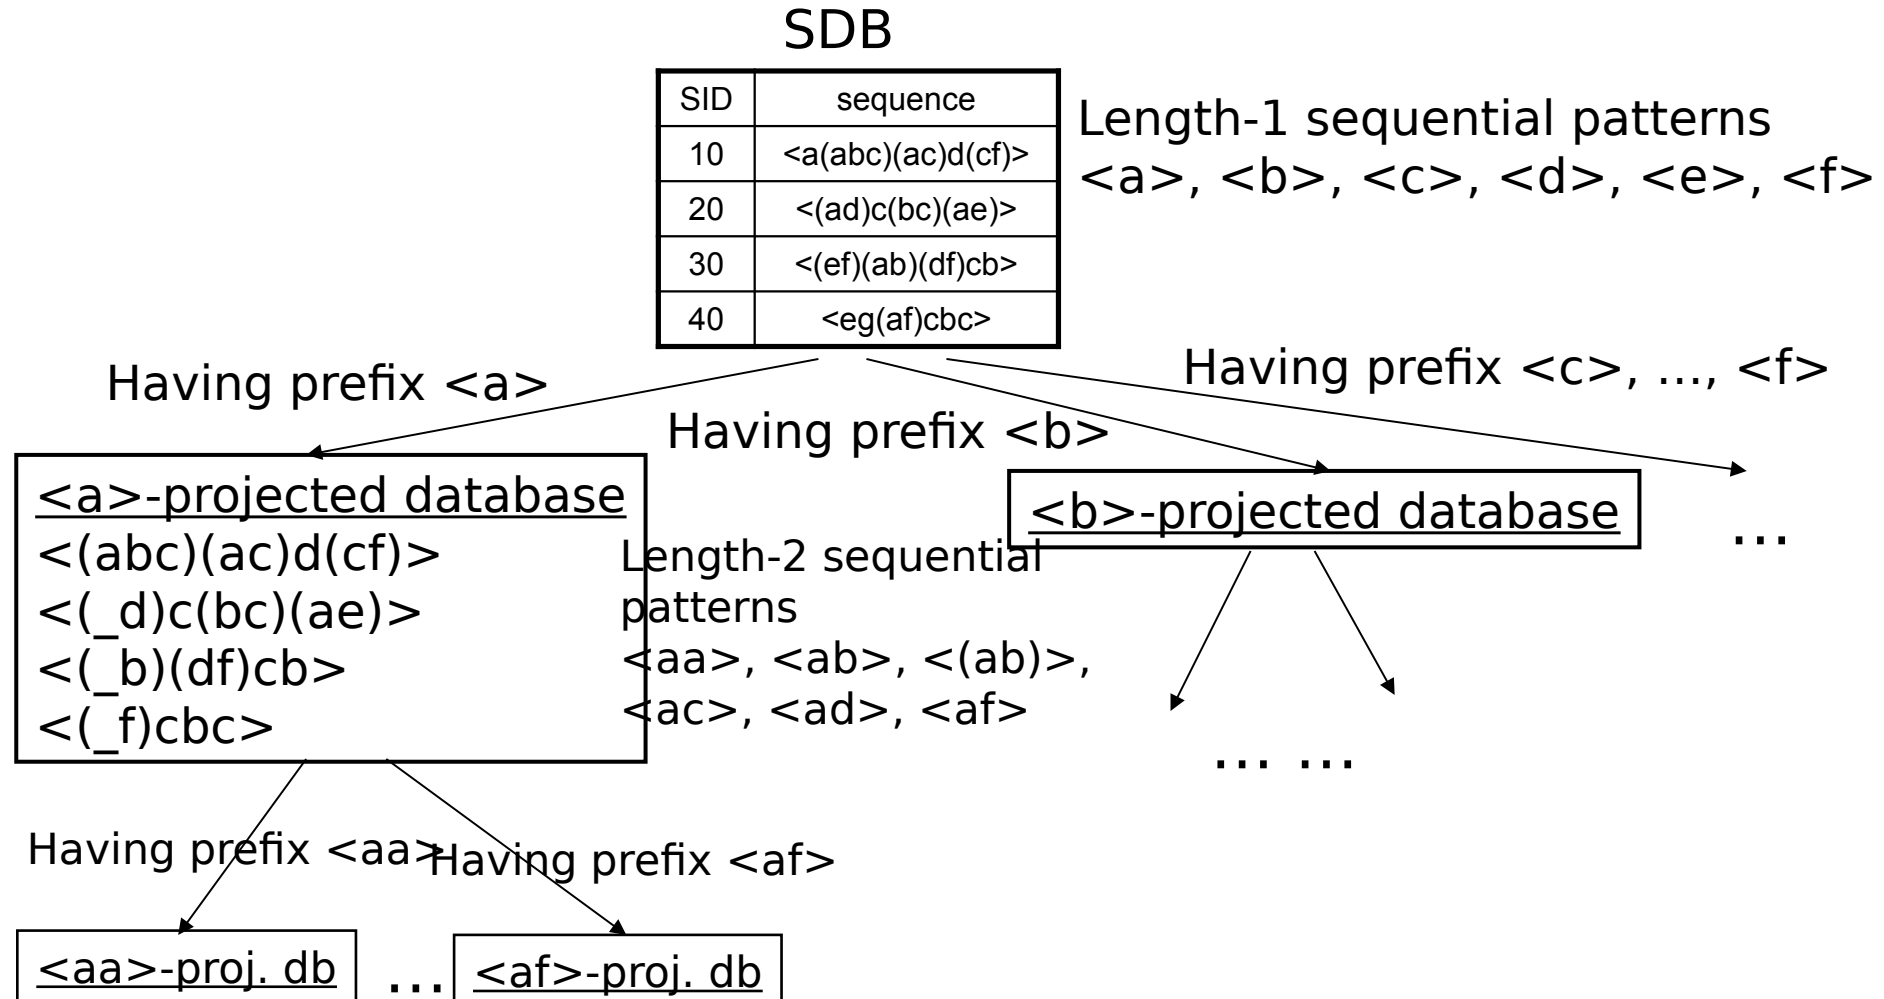
SID	sequence
10	$\langle a(abc)(ac)d(cf) \rangle$
20	$\langle (ad)c(bc)(ae) \rangle$
30	$\langle (ef)(ab)(df)cb \rangle$
40	$\langle eg(af)c bc \rangle$

Finding Sequence Patterns with Prefix <a>

- Only need to consider projections with respect to <a>
 - <a>-projected database: <(abc)(ac)d(cf)>, <(_d)c(bc)(ae)>, <(_b)(df)cb>, <(_f)cbc>
- Find all the length-2 seq. pat. Having prefix <a>: <aa>, <ab>, <(ac)>, <ad>, <ae>, <af>
 - Further partition into 6 subsets
 - Having prefix <aa>;
 - ...
 - Having prefix <af>

SID	sequence
10	<a(abc)(ac)d(cf)>
20	<(ad)c(bc)(ae)>
30	<(ef)(ab)(df)cb>
40	<eg(af)cbc>

Completeness of PrefixSpan



The Algorithm of PrefixSpan

- **Input:** A sequence database S , and the minimum support threshold min_sup
- **Output:** The complete set of sequential patterns
- **Method:** Call $\text{PrefixSpan}(<>, 0, S)$
- **Subroutine** $\text{PrefixSpan}(\alpha, l, S|\alpha)$
- **Parameters:**
 - α : sequential pattern,
 - l : the length of α ;
 - $S|\alpha$: the α -projected database, if $\alpha \neq <>$; otherwise; the sequence database S

The Algorithm of PrefixSpan

- **Method**

1. Scan $S|\alpha$ once, find the set of frequent items b such that:
 - a) b can be assembled to the last element of α to form a sequential pattern;
 - b) $\langle b \rangle$ can be appended to α to form a sequential pattern.
2. For each frequent item b , **append it to α** to form a sequential pattern α' , and output α' ;
3. For each α' , construct α' -projected database $S|\alpha'$, and call PrefixSpan(α' , $l+1$, $S|\alpha'$).

Efficiency of PrefixSpan

- No candidate sequence needs to be generated
- Projected databases keep shrinking
- Major cost of PrefixSpan: constructing projected databases
 - Can be improved by bi-level projections

Optimization in PrefixSpan


- Single level vs. bi-level projection
 - Bi-level projection with 3-way checking may reduce the number and size of projected databases
- Physical projection (disk) vs. pseudo-projection (memory)
 - Pseudo-projection may reduce the effort of projection when the projected database fits in main memory
- Parallel projection vs. partition projection
 - Partition projection may avoid the blowup of disk space

Scaling Up by Bi-Level Projection

- Partition search space based on length-2 sequential patterns
- Only form projected databases and pursue recursive mining over bi-level projected databases

Speed-up by Pseudo-projection


- Major cost of PrefixSpan: projection
 - Postfixes of sequences often appear repeatedly in recursive projected databases
- When (projected) database can be held in main memory, **use pointers to form projections**
 - Pointer to the sequence
 - Offset of the postfix




Pseudo-Projection vs. Physical Projection

- Pseudo-projection avoids physically copying postfixes
 - Efficient in running time and space when database can be held in main memory
- However, it is not efficient when database cannot fit in main memory
 - Disk-based random accessing is very costly
- Suggested Approach:
 - Integration of physical and pseudo-projection
 - Swapping to pseudo-projection when the data set fits in memory


Performance on Data Set C10T8S8I8



Performance on Data Set Gazelle




Effect of Pseudo-Projection

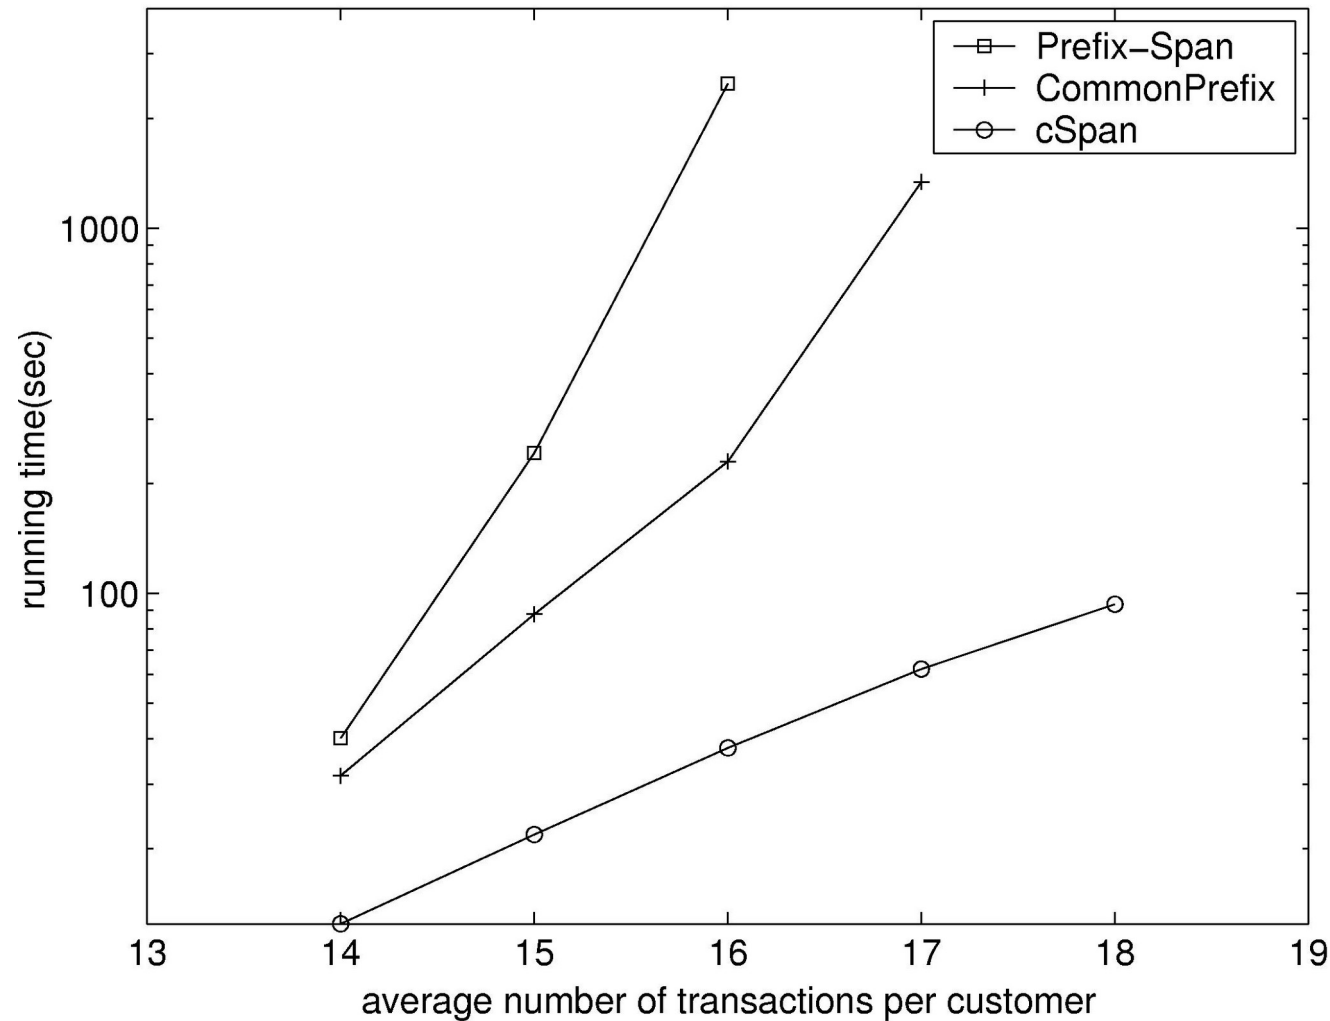


CloSpan: Mining Closed Sequential Patterns

- A **closed sequential pattern** s : there exists no superpattern s' such that $s' \supset s$, and s' and s have the same support
- **Motivation**: reduces the number of (redundant) patterns but attains the same expressive power
- Using Backward Subpattern and Backward Superpattern pruning to prune redundant search space



CloSpan: Performance Comparison with PrefixSpan



Extensions to Frequent Sequence Mining

Constraints for Seq.-Pattern Mining

- Item constraint
 - Find web log patterns only about online-bookstores
- Length constraint
 - Find patterns having at least 20 items

→
- Super pattern constraint
 - Find super patterns of “PC & digital camera”
- Aggregate constraint
 - Find patterns that the average price of items is over \$100

More Constraints

- Regular expression constraint
 - Find patterns “starting from Yahoo homepage, search for hotels in Washington DC area”
 - `Yahootravel(WashingtonDC|DC)(hotel|motel|lodging)`
- Duration constraint
 - Find patterns about ± 24 hours of an event
- Gap constraint
 - Find purchasing patterns such that “the gap between each consecutive purchases is less than 1 month”

From Sequential Patterns to Structured Patterns

- Sets, sequences, trees, graphs, and other structures
 - Transaction DB: Sets of items
 - $\{\{i_1, i_2, \dots, i_m\}, \dots\}$
 - Seq. DB: Sequences of sets:
 - $\{\langle\{i_1, i_2\}, \dots, \{i_m, i_n, i_k\}\rangle, \dots\}$
 - Sets of Sequences:
 - $\{\{\langle i_1, i_2 \rangle, \dots, \langle i_m, i_n, i_k \rangle\}, \dots\}$
 - Sets of trees: $\{t_1, t_2, \dots, t_n\}$
 - Sets of graphs (mining for frequent subgraphs):
 - $\{g_1, g_2, \dots, g_n\}$
- Mining structured patterns in XML documents, bio-chemical structures, etc.

Episodes and Episode Pattern Mining

- Other methods for specifying the kinds of patterns
 - Serial episodes: $A \rightarrow B$
 - Parallel episodes: $A \ \& \ B$
 - Regular expressions: $(A \mid B)C^*(D \rightarrow E)$
- Methods for episode pattern mining
 - Variations of Apriori-like algorithms, e.g., GSP
 - Database projection-based pattern growth
 - Similar to the frequent pattern growth without candidate generation

Periodicity Analysis

- Periodicity is everywhere: tides, seasons, daily power consumption, etc.
- **Full periodicity**
 - Every point in time contributes (precisely or approximately) to the periodicity
- **Partial periodicity:** A more general notion
 - Only some segments contribute to the periodicity
 - Jim reads NY Times 7:00-7:30 am every week day
- **Cyclic association rules**
 - Associations which form cycles
- **Methods**
 - Full periodicity: FFT, other statistical analysis methods
 - Partial and cyclic periodicity: Variations of Apriori-like mining methods

Summary

- Sequential Pattern Mining is useful in many application, e.g. weblog analysis, financial market prediction, BioInformatics, etc.
- It is similar to the frequent itemsets mining, *but* with consideration of ordering.
- We have looked at different approaches that are descendants from two popular algorithms in mining frequent itemsets
 - Candidates Generation: AprioriAll and GSP
 - Pattern Growth: FreeSpan and PrefixSpan