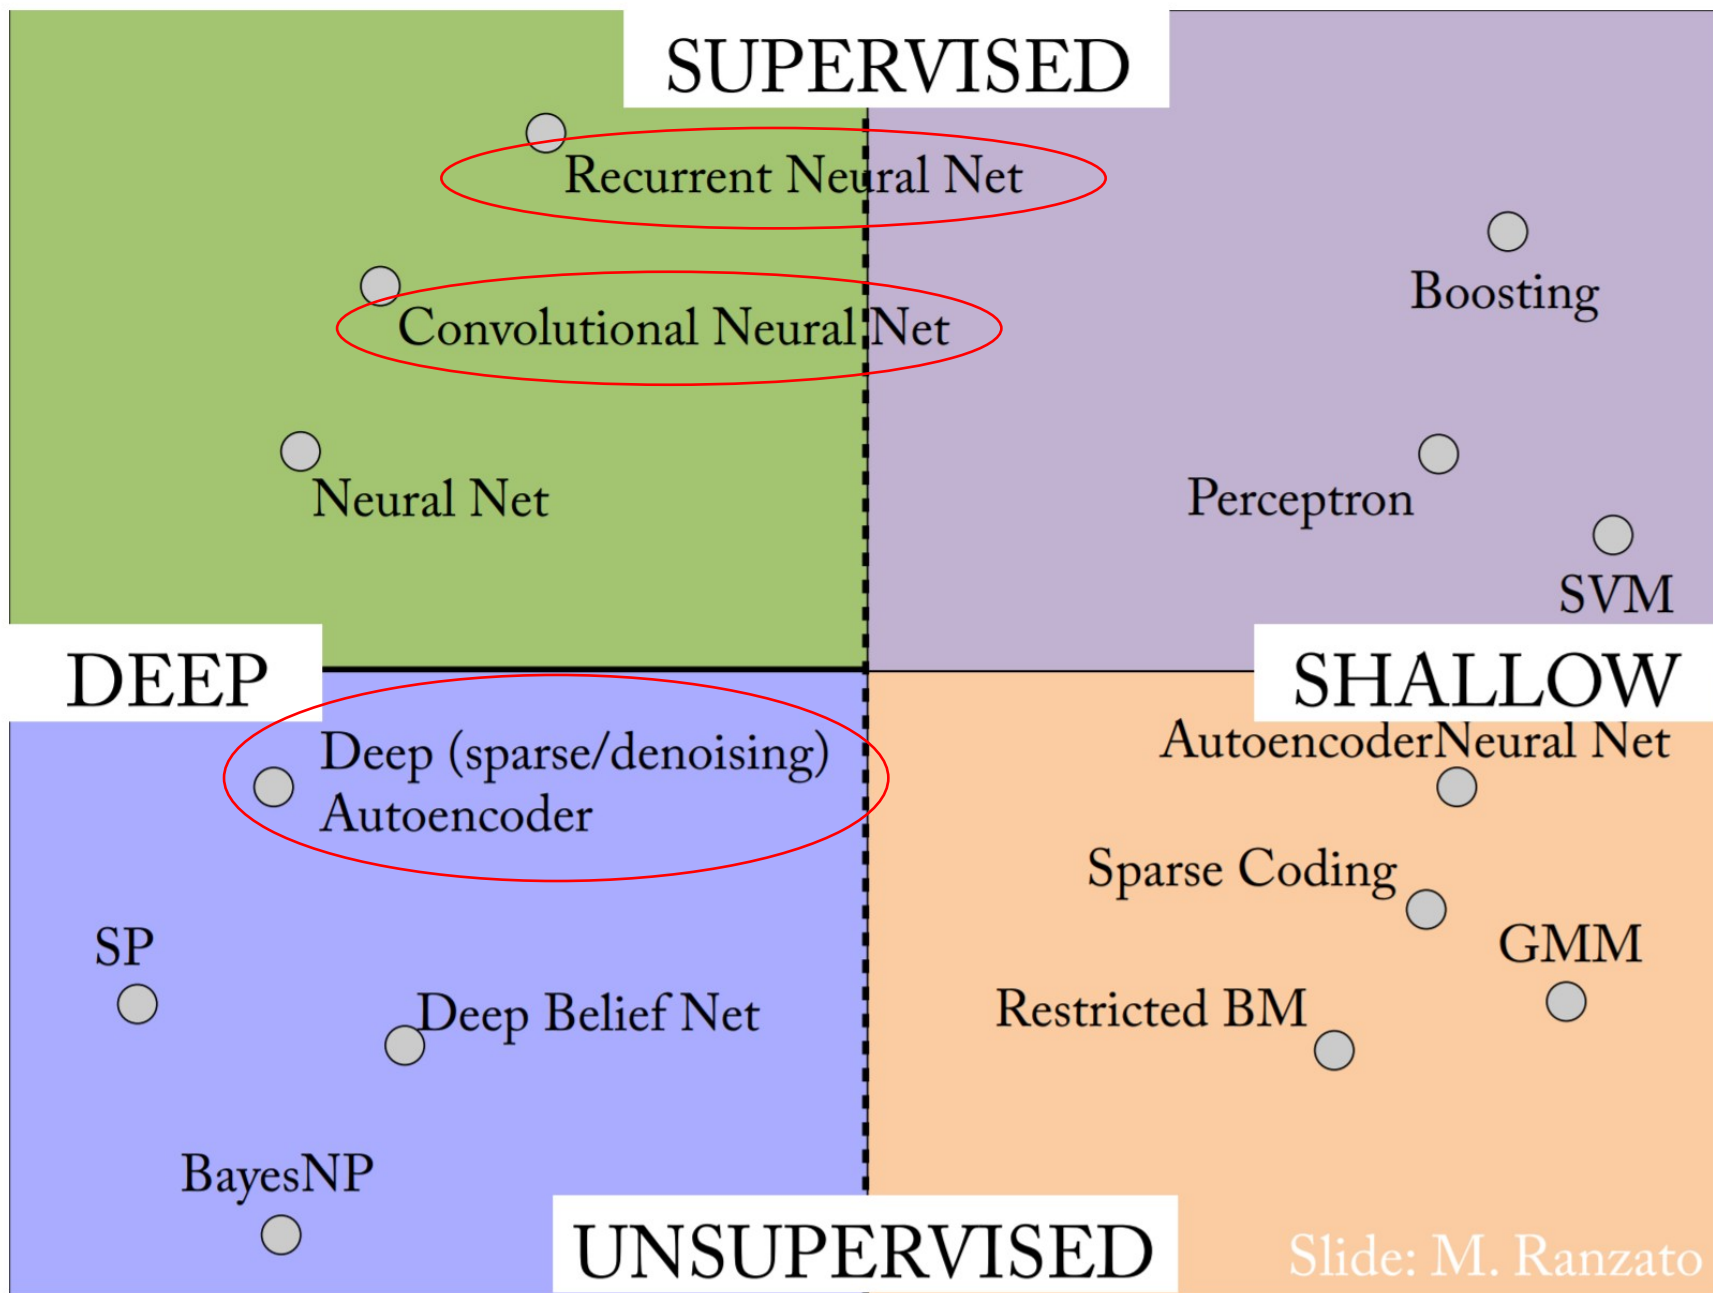


Deep Learning

Deep Learning

- Most conventional neural networks (MLP) have a few layers
- Deep networks have **many** layers
- Multiple layers can build an improved feature space
 - First layer learns first order features (edges in an image)
 - Second layer learns higher order features (combinations of edges)
 - And so on...
 - Final layer of transformed features is fed into a set of fully connected layers with backpropagation.



Example: Classifying Hand Written Digits



Example: Image Captioning



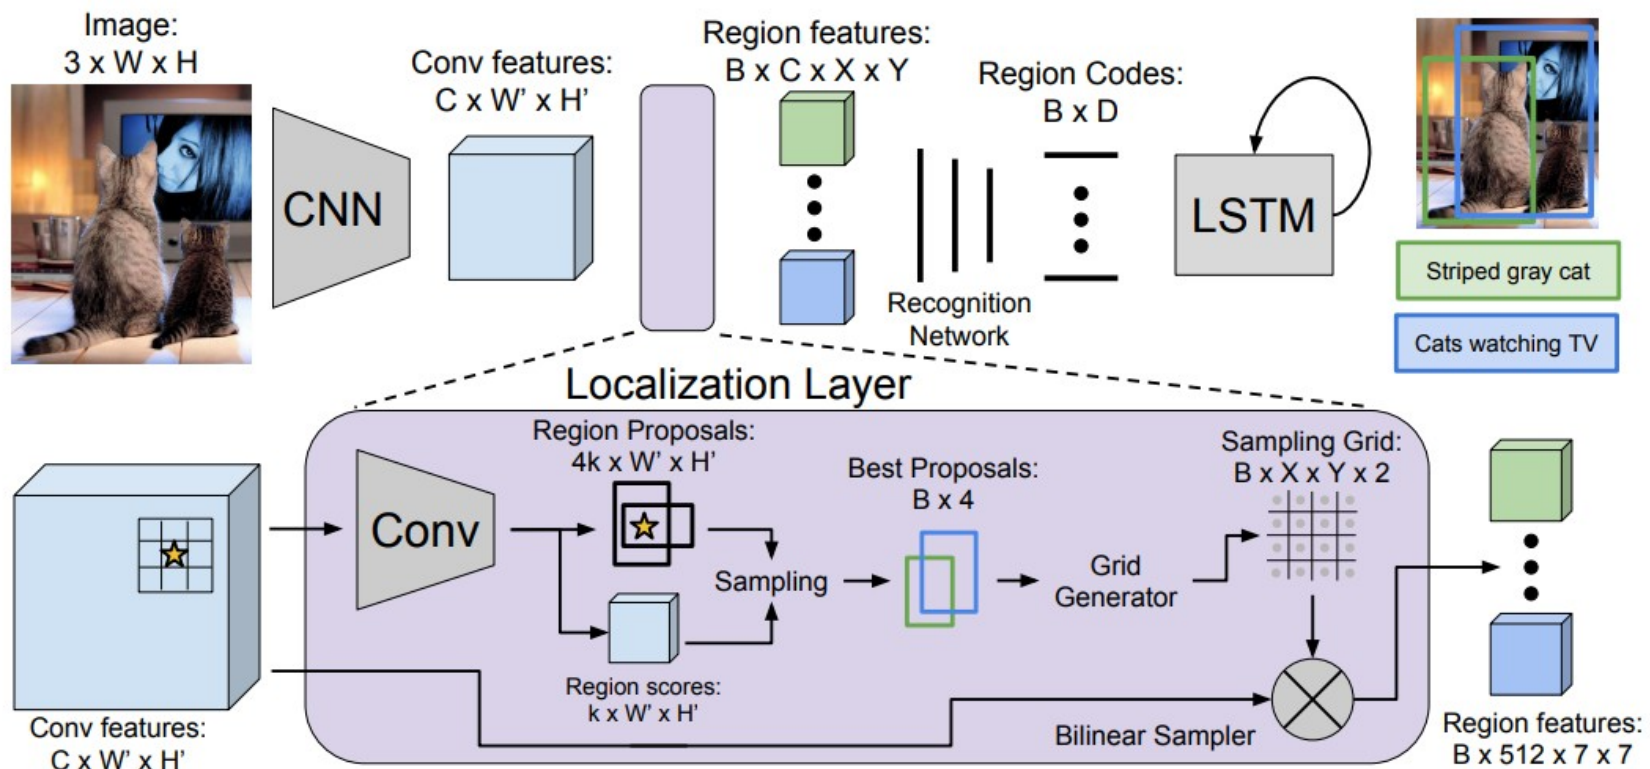
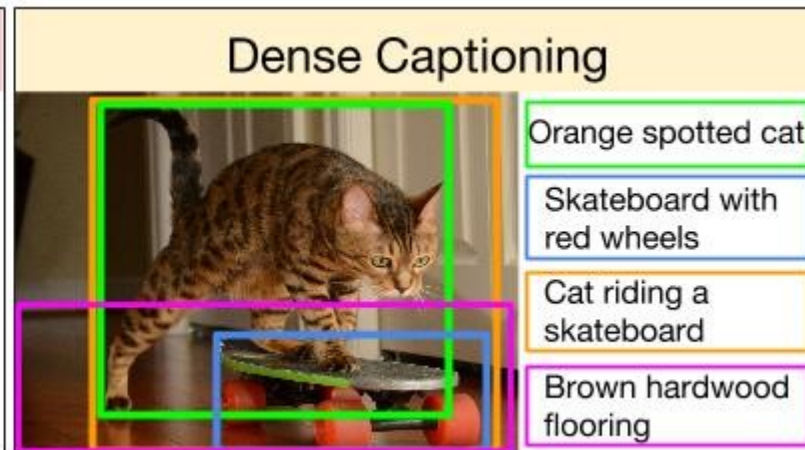
a soccer player is kicking a soccer ball



a street sign on a pole in front of a building



a couple of giraffe standing next to each other




DenseCap (Johnson et al., 2016)


Convolutional Neural Networks (CNNs)

- Emerged from the study of the visual cortex of cats
- Based on the idea of a *local receptive field*
- The local receptive fields of different neurons may overlap
- Some neurons react only to a particular stimuli
- Some neurons react to more complex patterns that are combinations of lower level patterns
- Higher level neurons are based on the outputs of neighboring lower level neurons

The Primary Visual Pathway





A Simple CNN



Back to the Data

- CNNs are typically used to classify **images**
- Images are a matrix of pixel values
- Grayscale images have one channel
- Color Images have 3 channels (red, green, and blue)



Convolution Step

- The purpose of convolution is to **extract features** from the image
- The convolution step in a CNN applies a sliding window-like filter over the image
- For every position, the element wise multiplication between the two matrices is computed to create the convolved feature

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

1	0	1
0	1	0
1	0	1

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0





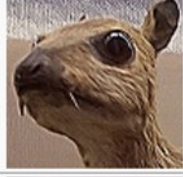
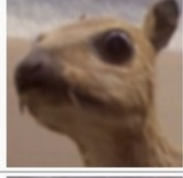
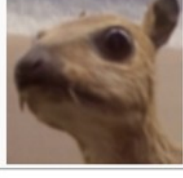
Image

4		

Convolved
Feature

Types of Filters

- There are different types of filters that can be applied to an image for feature detection.
- Different filters can detect different features from an image

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

Another Example




Input

Considerations for Convolution

- **Depth:** corresponds to the number of filters we use for the convolution operation.
 - Different filters find different features
 - Three filters result in a feature map depth of three
- **Stride:** the number of pixels by which we slide our filter matrix over the input matrix.
 - When the stride is 1 then we move the filters one pixel at a time.
 - When the stride is 2, then the filters jump 2 pixels at a time as we slide them around.
 - Having a larger stride will produce smaller feature maps.
- **Zero-padding:** padding the input matrix with zeros around the border
 - Allows the filter to be applied to bordering elements of our input image matrix.
 - Allows for the control of the size of the feature maps.

Introducing Non-Linearity

- Convolution is a linear operation
- Rectified linear unit function (ReLU)
- Applied per pixel of the feature map
- Computes a linear function of the inputs and outputs the result if it is positive and 0 otherwise
Output = $\text{Max}(0, \text{Input})$
- Effectively replaces all negative values in the feature map with 0
- Other functions
 - tanh, sigmoid, softmax



ReLU Example

Input Feature Map



ReLU

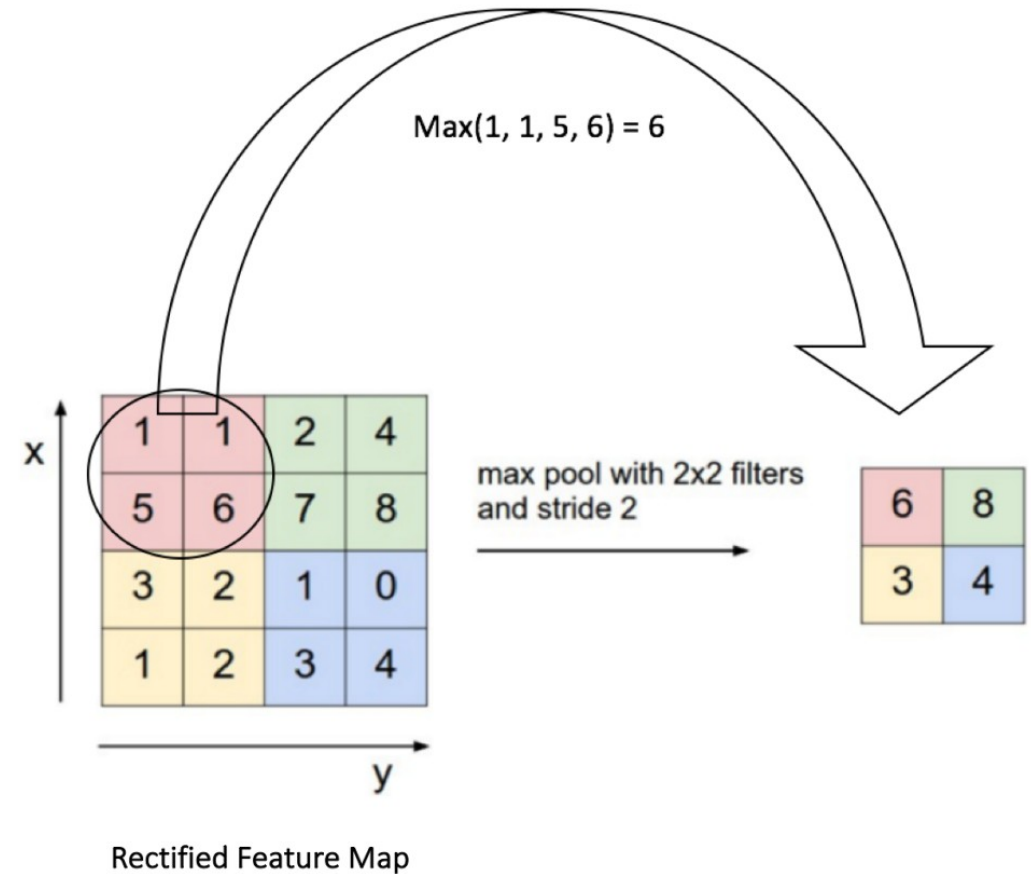


Rectified Feature Map



Pooling Step

- **Subsamples** the feature map
- Reduces the size of each feature map retaining important information
- Can be of different types (max, average, sum)
- Parameters
 - size of kernel
 - Stride
 - padding




Pooling Example




Rectified Feature Map


Pooling



Max




Sum

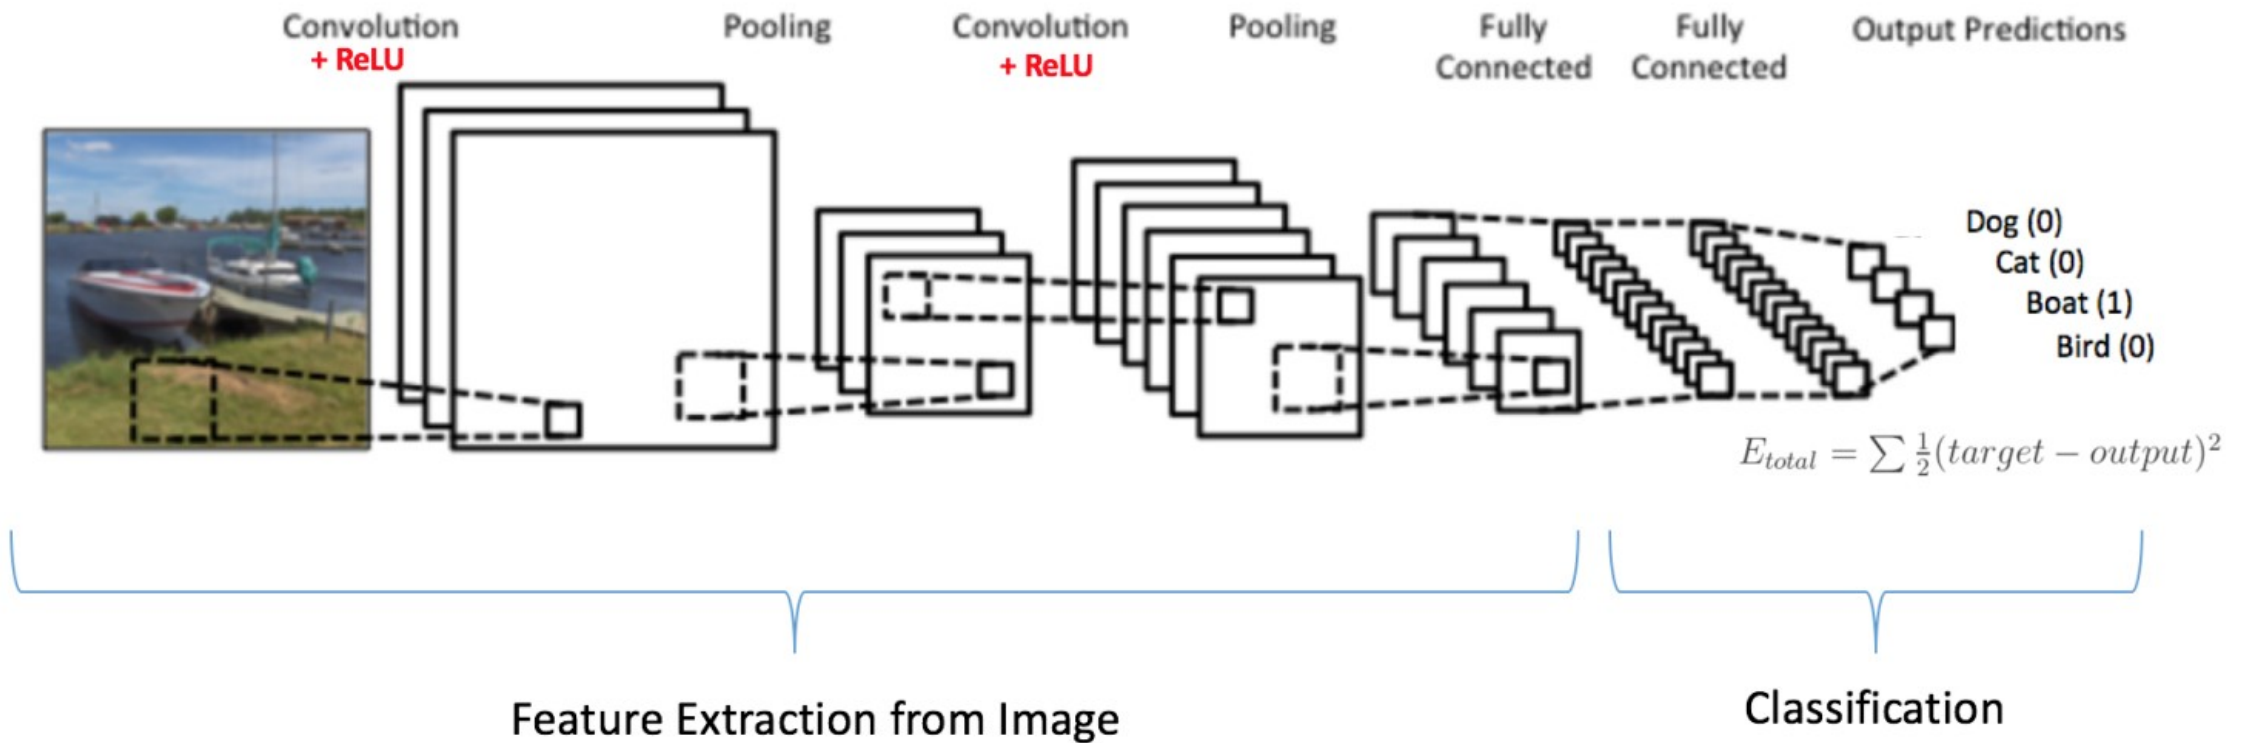


Fully Connected Layers

- Features of pooling layer is used as input to fully connected layer
- Output of convolution and pooling layers represent high-level features of the image
- Fully connected layer is a traditional multilayer perceptron.
- High level features are fed to the fully connected layer to classify the image




Training the Network




The Problem with HyperParameters

- CNNs have a large number of hyperparameters
 - Number of layers
 - Convolution parameters (depth, stride, padding)
 - Non-linear function (ReLU, softmax, sigmoid, tanh)
 - Pooling layer parameters (kernel size, stride, padding)
 - Fully connected layer (number of hidden layers, optimization algorithm, learning rate, momentum term, etc...)
 - Weight initialization
 - Dropout (reduces overfitting)
- Still a lot of open issues

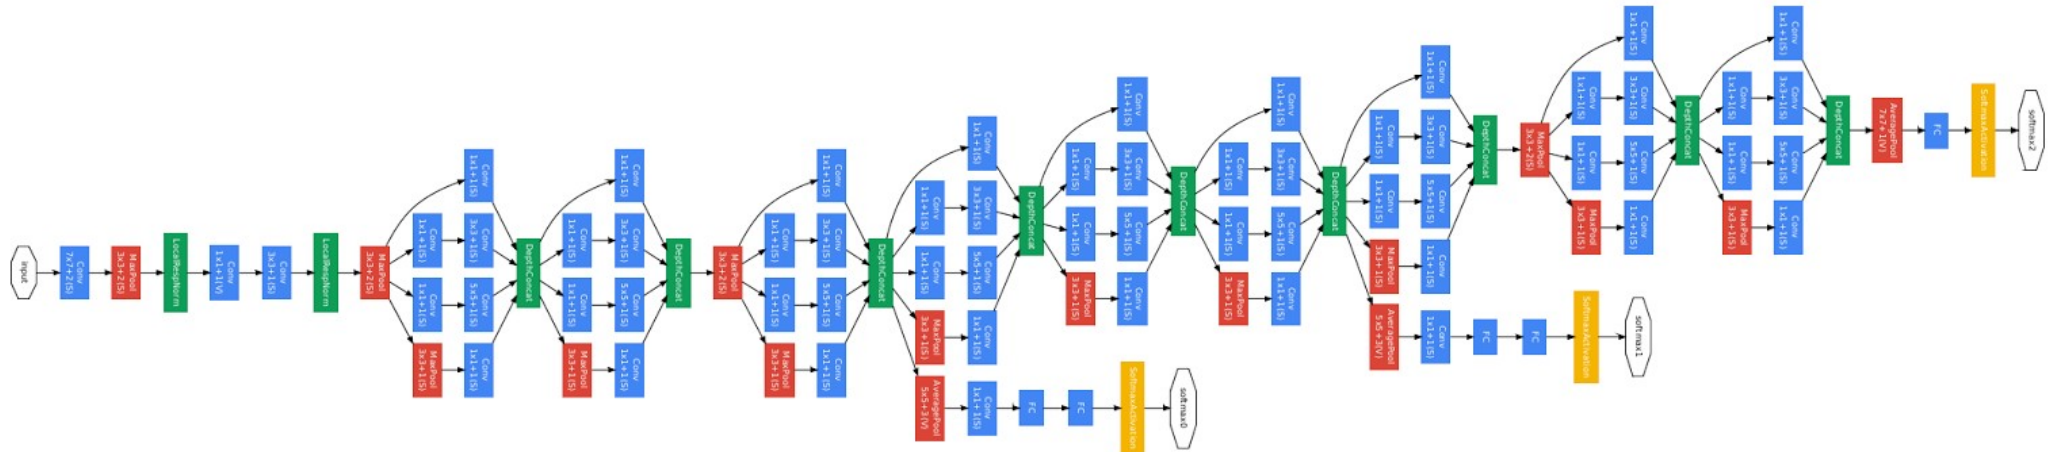
Popular CNN Architectures: LeNet-5 (LeCun, 1998)



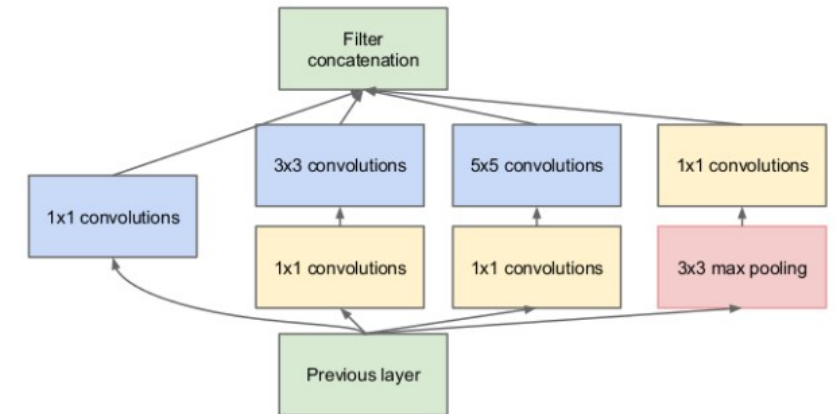
Popular CNN Architectures: AlexNet (Krizhevsky, 2012)




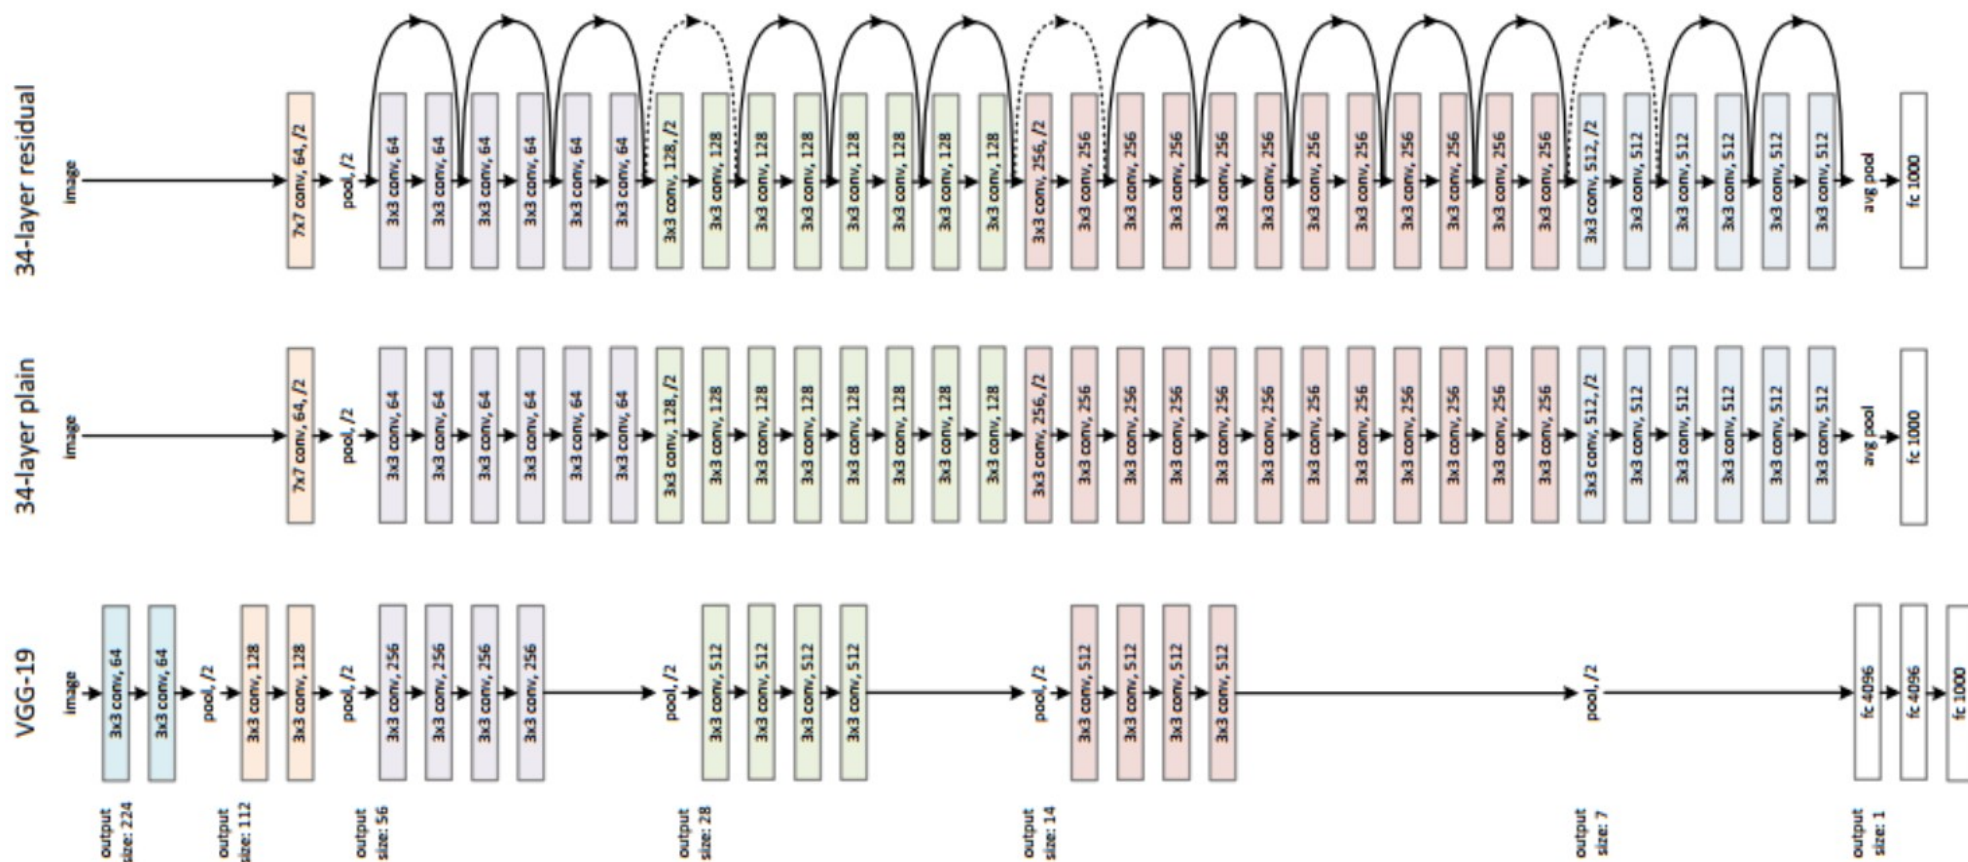
Popular CNN Architectures: GoogLeNet (Szegedy et al., 2015)




Inception Module



Popular CNN Architectures: ResNet (He, 2015)



CNN Demo



Recurrent Neural Networks

- *We are the sum total of our experiences. None of us are the same as we were yesterday, nor will be tomorrow. (B.J. Neblett)*

Recurrent Neural Networks

- Idea: our brain learns from experience and that experience is used to solve new problems
- Human learning can be distilled into the following:
 - Memorization: every time we gain new information, we store it for future reference.
 - Combination: not all tasks are the same, so we couple our analytical skills with a combination of our memorized, previous experiences to reason about the world.
- Our brains work with sequences to learn

Recurrent Neural Networks

- Traditional neural networks operate on the fundamental assumption that inputs are independent.
 - The output at any given time is completely determined by the input at the same time.
- Recurrent neural networks map an entire history of inputs to each output.
- Can be thought of as a traditional neural network with a loop

Memory in Recurrent Neural Networks

Weights in a Recurrent
Neural Network

$$y_0 = f(W_x X_0)$$

$$y_1 = f(W_x X_1)$$

$$y_2 = f(W_x X_2)$$

Weights with Current Input
and Previous Input

$$y_0 = f(W_x X_0)$$

$$y_1 = f(W_x X_1 + W_h X_0)$$

$$y_2 = f(W_x X_2 + W_h X_1)$$


Weights with a
Combination of Current
Input and All Previous
Inputs

$$y_0 = f(W_x X_0)$$


$$y_1 = f(W_x X_1 + W_h f(W_x X_0))$$

$$y_2 = f\left(W_x X_2 + W_h f(W_x X_1 + W_h f(W_x X_0))\right)$$

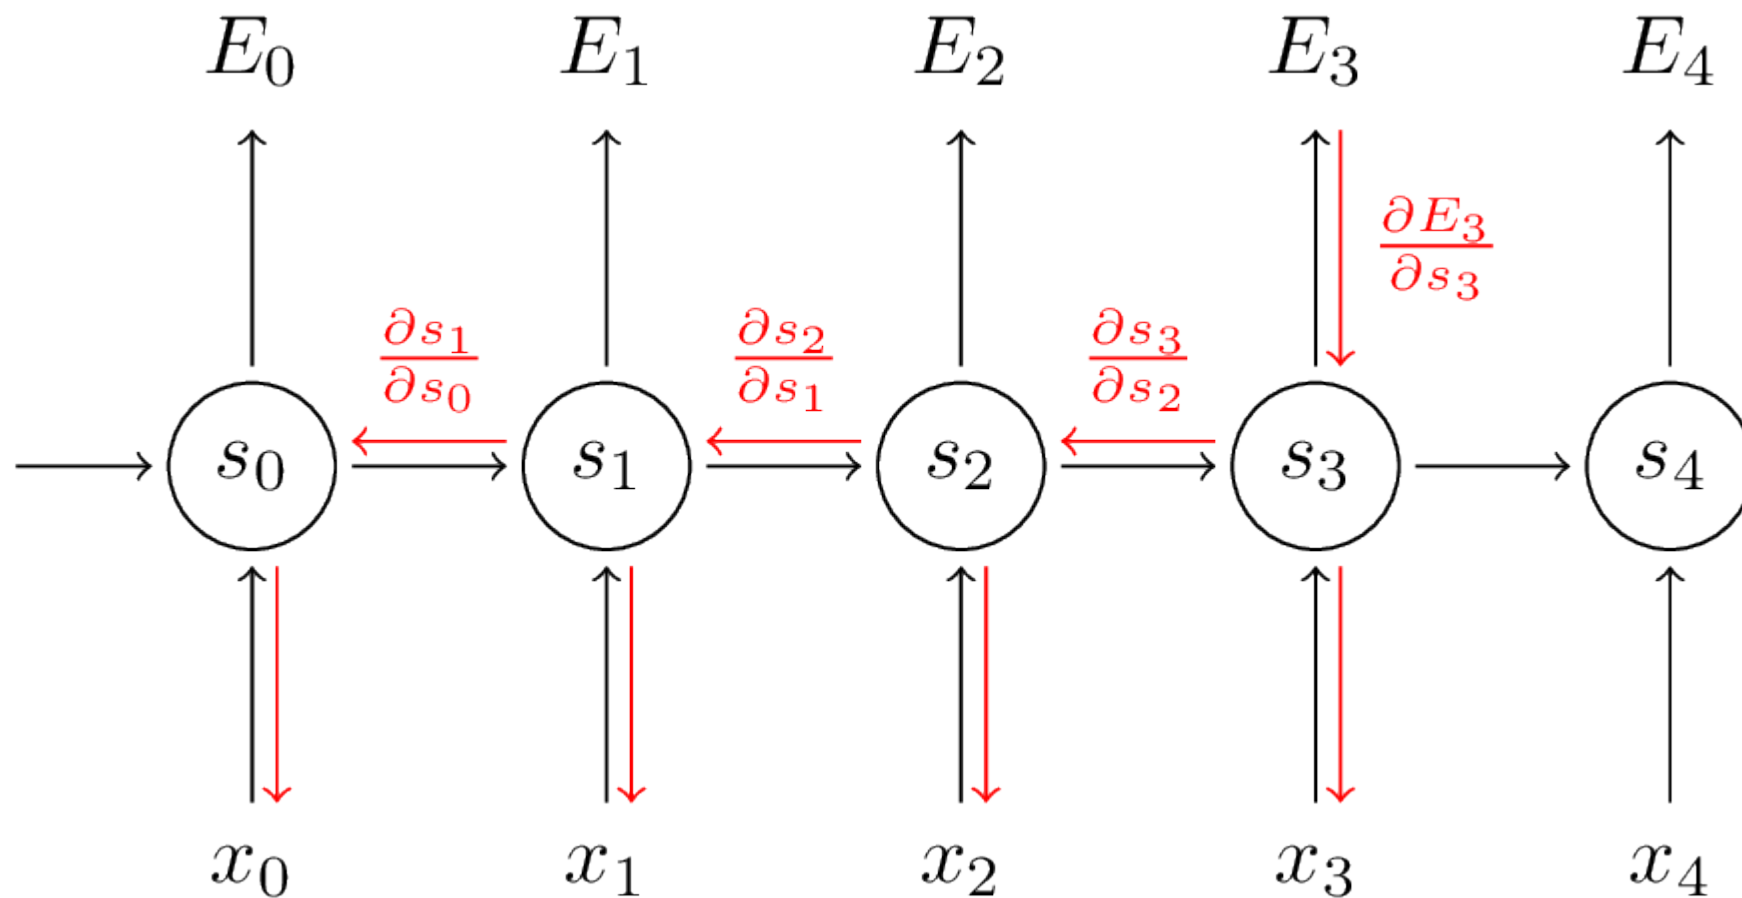
An Unrolled Recurrent Neural Network



Standard RNN Architecture




Back Propagation in an RNN

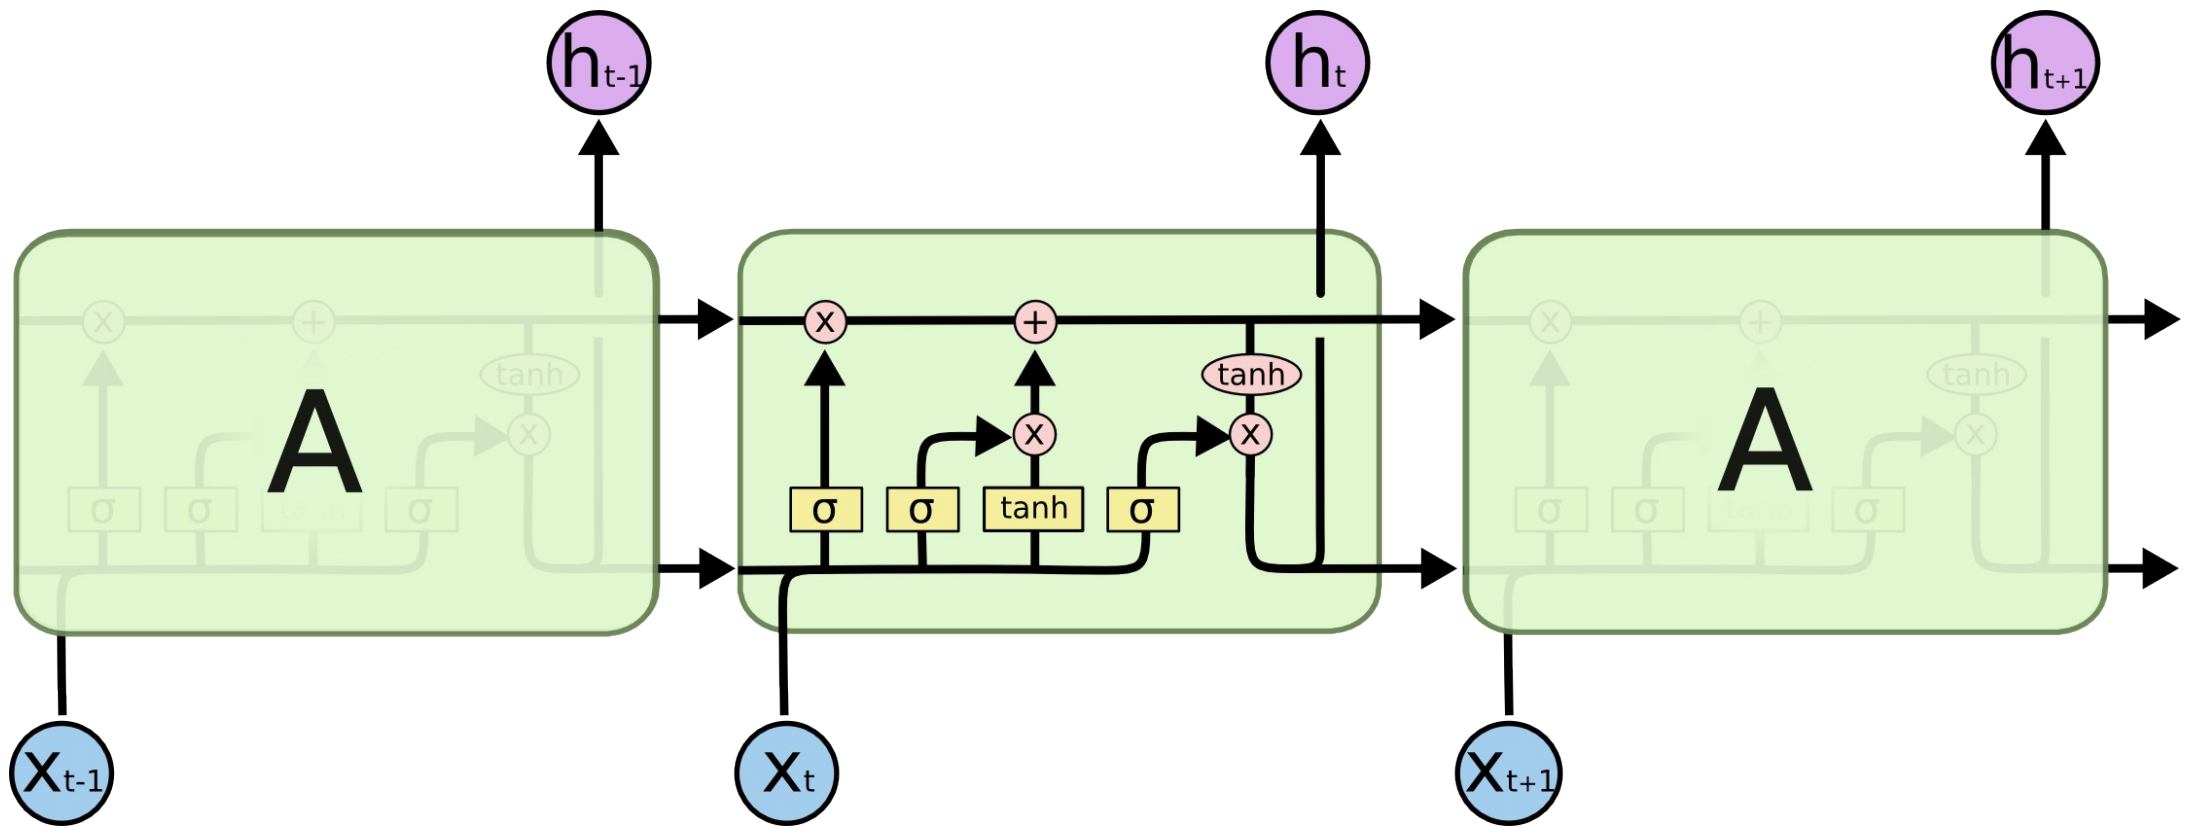


The Vanishing Gradient Problem

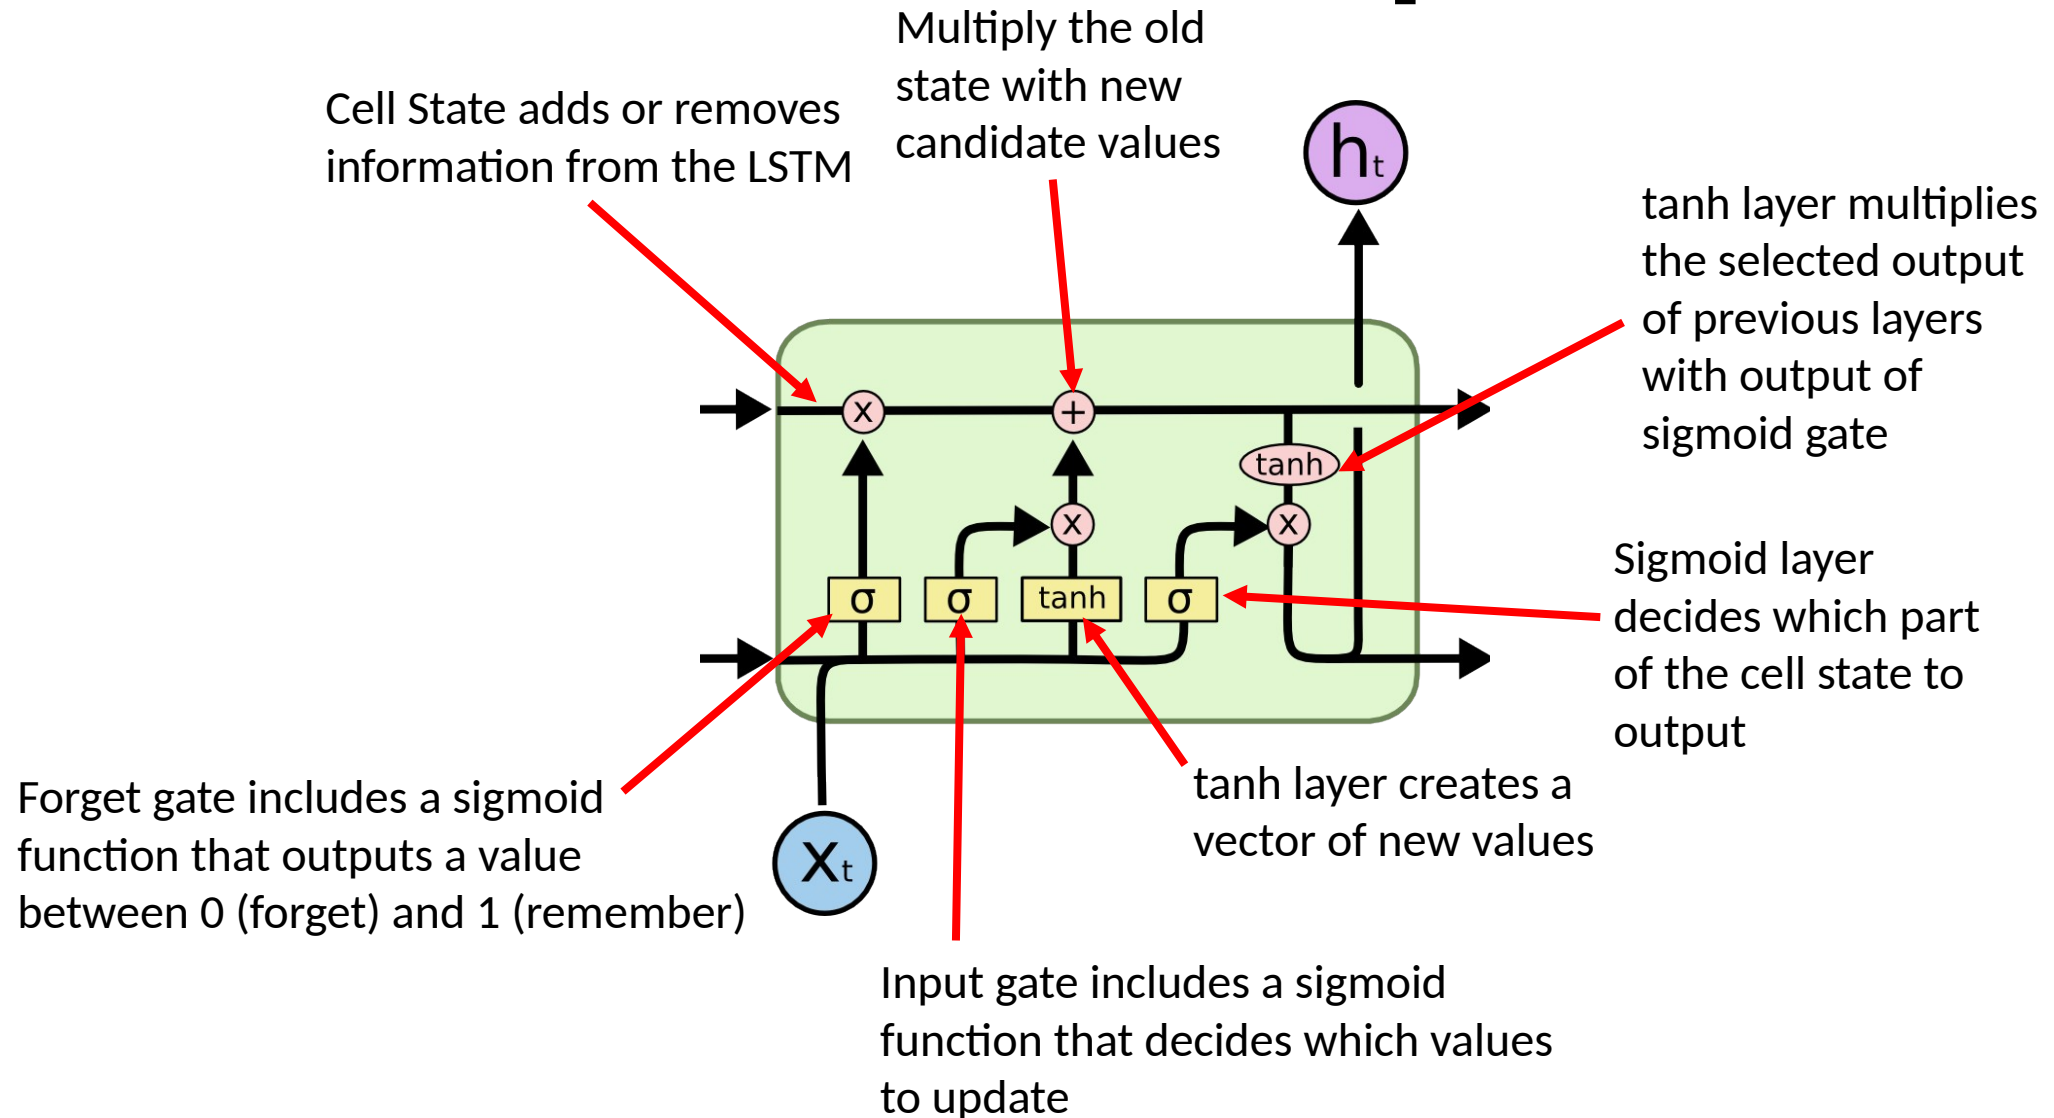
- RNNs have a difficult time with long term dependencies
- Activation functions such as tanh and sigmoid have derivatives of 0 at both ends.
- Gradient contributions from “far away” steps become zero, and the state at those steps doesn’t contribute to what you are learning



LSTM RNNs



LSTM RNN Node Components



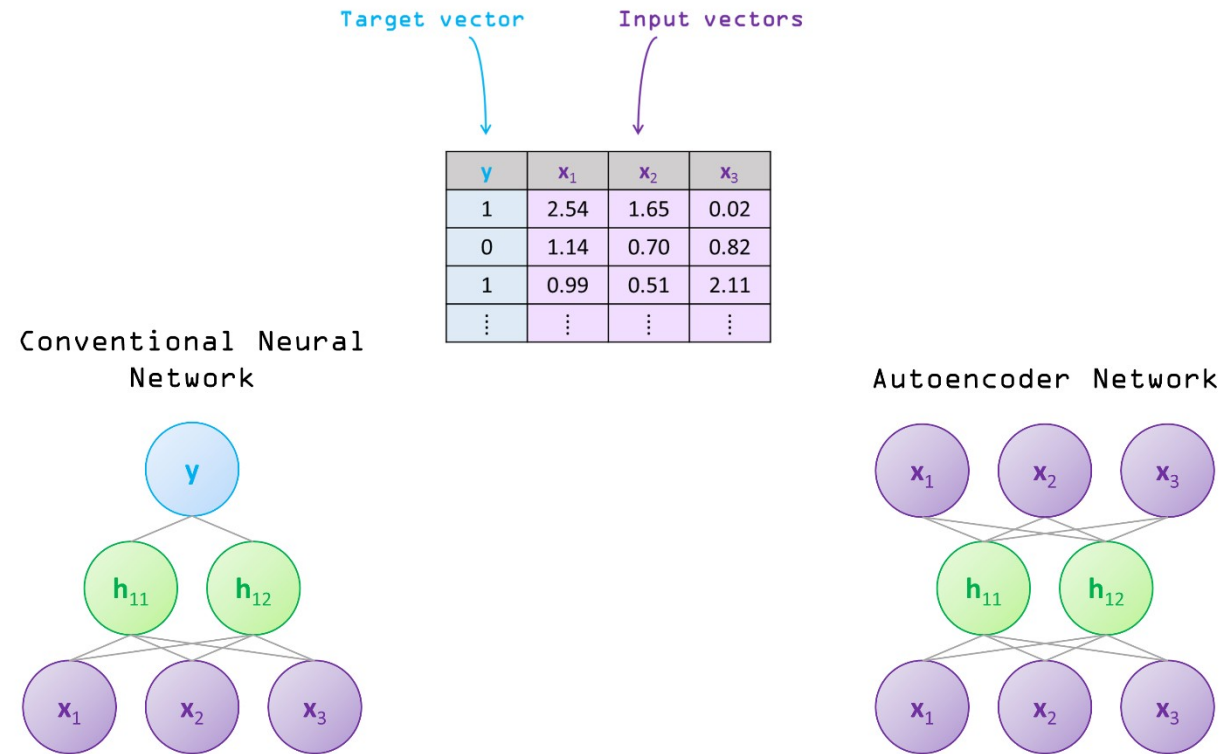
RNN Handwriting Generation Demo

- <http://www.cs.toronto.edu/~graves/handwriting.html>


Autoencoders for Unsupervised Feature Detection

- What about when we don't have labeled data?
- Neural networks can be used to learn efficient representations or codings of the input data.
- Codings have a much lower dimensionality than input data
- Autoencoders are powerful feature detectors
- Used for pretraining deep neural networks
- Used to generate new data that looks like the input

Autoencoders






Autoencoders



Denoising Autoencoders

- Denoising autoencoders can be used to learn useful features by adding noise to its inputs







Reconstructed Features and Learned Features

7



2




Stacked Denoising Autoencoders

- Uses stacks of autoencoders to learn hidden representations of the data
- As we saw before, denoising autoencoders learn from the data by adding noise and trying to reconstruct it
- Each layer consists of a denoising autoencoder with input from the previous layer's output
- Trained to minimize reconstruction error in the output

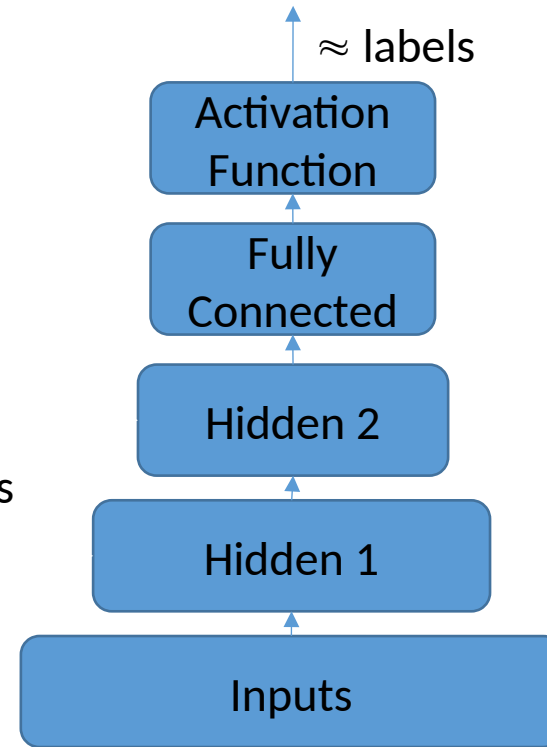
Autoencoders for Unsupervised Pretraining

- Sometimes you have a supervised learning task but limited label data



Phase 1: Train the autoencoder using all the data

Copy Parameters



Phase 1: Train the autoencoder using all the data